Master thesis

Master in Research and Innovation

Inferring programs structure from an execution trace

Author

Juan Francisco Martínez Vera

Supervisor

Jesús Labarta Mancho

Facultat d'Informàtica de Barcelona (FIB)
Universitat Politècnica de Catalunya (UPC)

Outline

Performance Analysis

- Aid to detect **bottlenecks**
  - That prevents from better performance
- Requires high skilled analyst...
- ... so developers are not used to work with them
  - But derive this work to actual specialist
  - Analyst are used to **work with codes they are not familiar with**.
  - e.g. POP project: Provides **performance optimisation and productivity** services for academic and industrial code(s).

Motivations

About improving understandability

Providing application structure will lead to better understanding about what the application is doing.

Motivations

About improving understandability

Providing application structure will lead to better understanding about what the application is doing.

About improving reports

Having the structure of the application, the communication between analyst and developer can be improved.

Objective

The objective for this thesis is...

Expose the structure of the application, by means of a post-mortem trace analysis.

State of the Art

First step is always to explore what is over there...

- Some previous research have been driven in this field.
- Since traces are ordered sequences...
- ... Natural choice has been **sequential pattern mining** in general.

State of the Art

First step is always to explore what is over there...

- Some previous research have been driven in this field.
- Since traces are ordered sequences...
- ... Natural choice has been **sequential pattern mining** in general.

Just **pick and develop one of the proposals?**

State of the Art

First step is always to explore what is over there...

- Some previous research have been driven in this field.
- Since traces are ordered sequences...
- ... Natural choice has been **sequential pattern mining** in general.

Just **pick and develop one of the proposals?**

Looking to the trend of increasing trace sizes...

- This sort of algorithms, **used to present high complexity**.
- From $O(n^2)$ to $O(2^n)$.

State of the Art

First step is always to explore what is over there...

- Some previous research have been driven in this field.
- Since traces are ordered sequences...
- ... Natural choice has been **sequential pattern mining** in general.

Just **pick and develop one of the proposals?**

Looking to the trend of increasing trace sizes...

- This sort of algorithms, **used to present high complexity**.
- From $O(n^2)$ to $O(2^n)$.

Finally...

The decision has been to explore a new approach!

Outline

Application structure by classification

The key idea

Use communications as **proxies** for the observation of iterations and clustering them **by its behaviour**.

Application structure by classification

The key idea

Use communications as **proxies** for the observation of iterations and clustering them **by its behaviour**.

Why?

HPC applications idiosincracy

- Big outer loop
- Repetitive and stable executions
- Communications (MPI) lies on loops that drives the execution

So...

- **Similar behaviour for all iterations** in a given loop
- Loops by **monitoring the communications**.

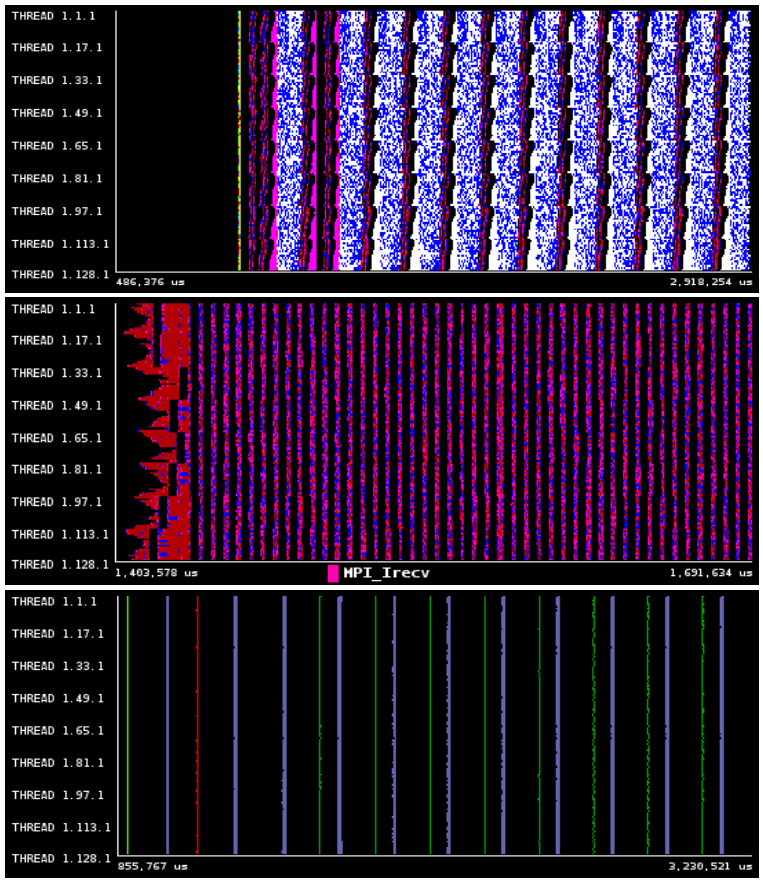

Application structure by classification

What will define **behaviour**?
Selected features must be able to

- Join MPIs from the same loop
- Separe MPIs from different loops

Application structure by classification

What will define **behaviour**?
Selected features must be able to

- Join MPIs from the same loop
- Separate MPIs from different loops

As a starting point ...

1. **Number of repetitions**: Two different mpi calls in same loop will be executed the same number of times.
2. **Mean time between repetitions**: Two different loops will, presumibly, execute different work.

Workflow

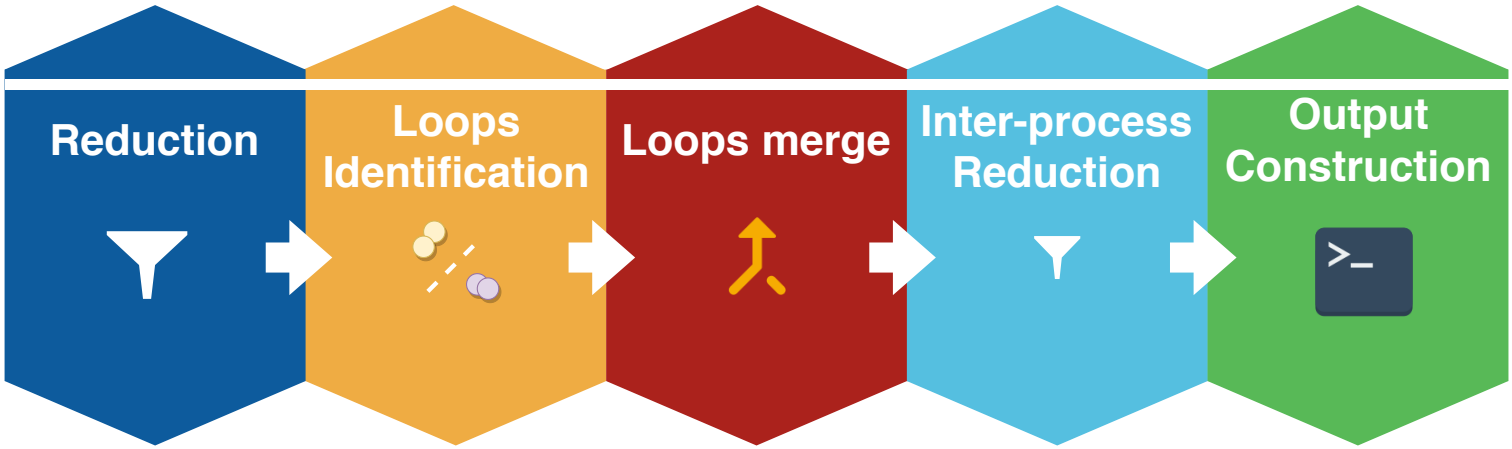
Reduction
Loops Identification
Loops merge
Inter-process Reduction
Output Construction
>_

Workflow

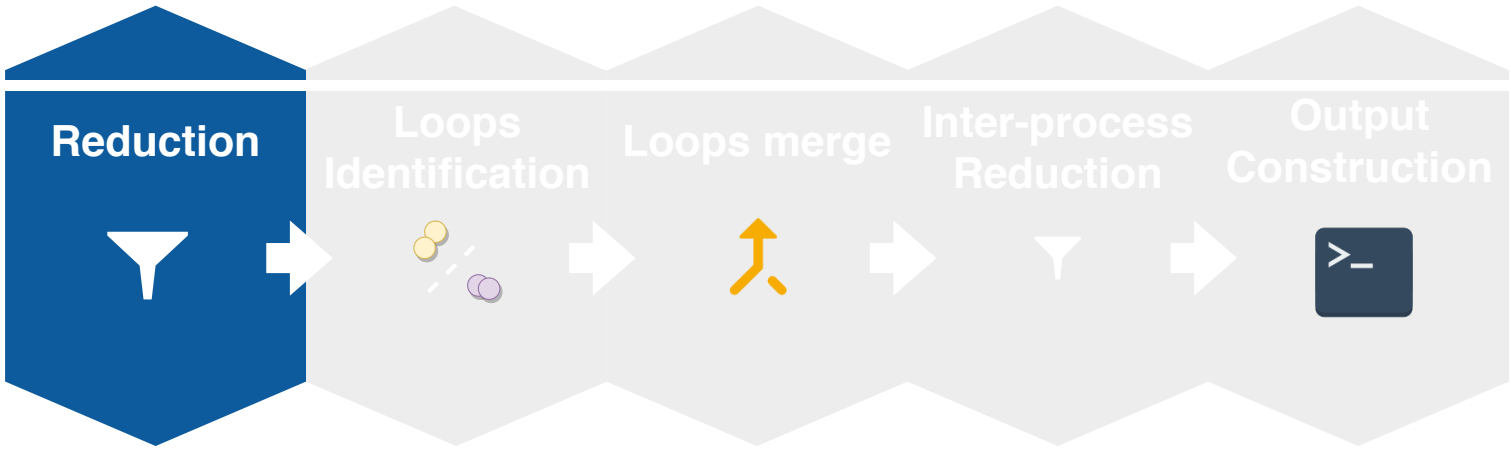

- Parsing and reducing information from trace

Workflow

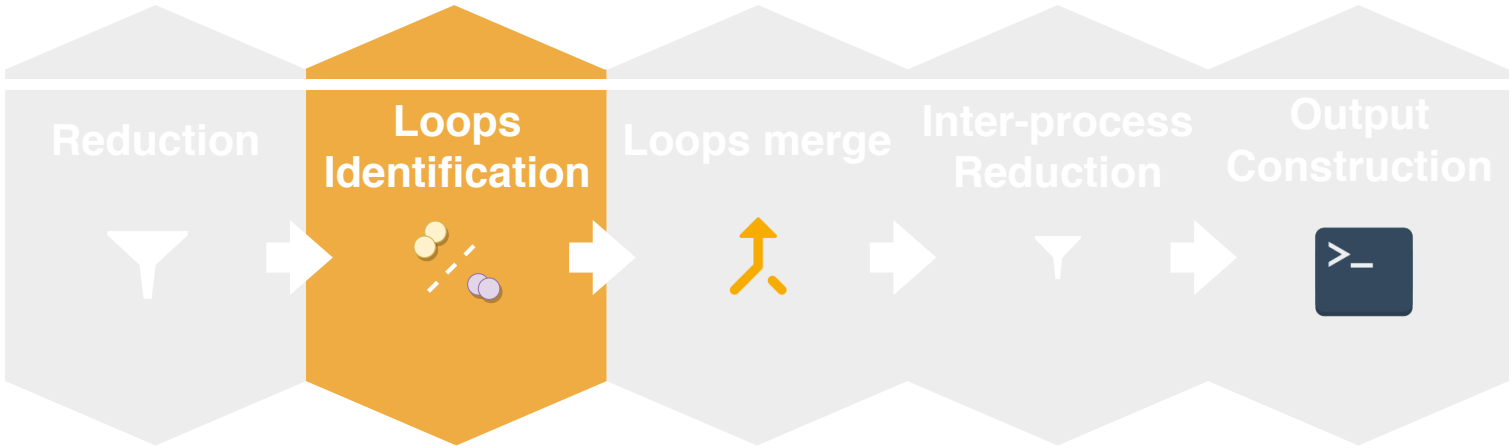

- Identify loops by MPI call sites classification

Workflow

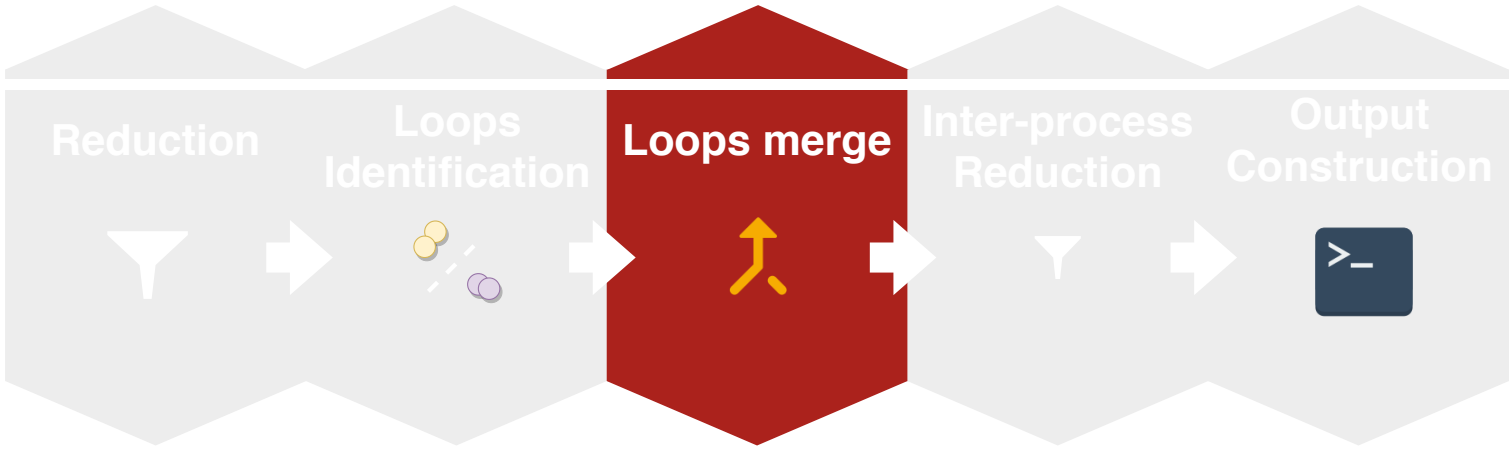

- Find out the hierarchical relations between loops

Workflow

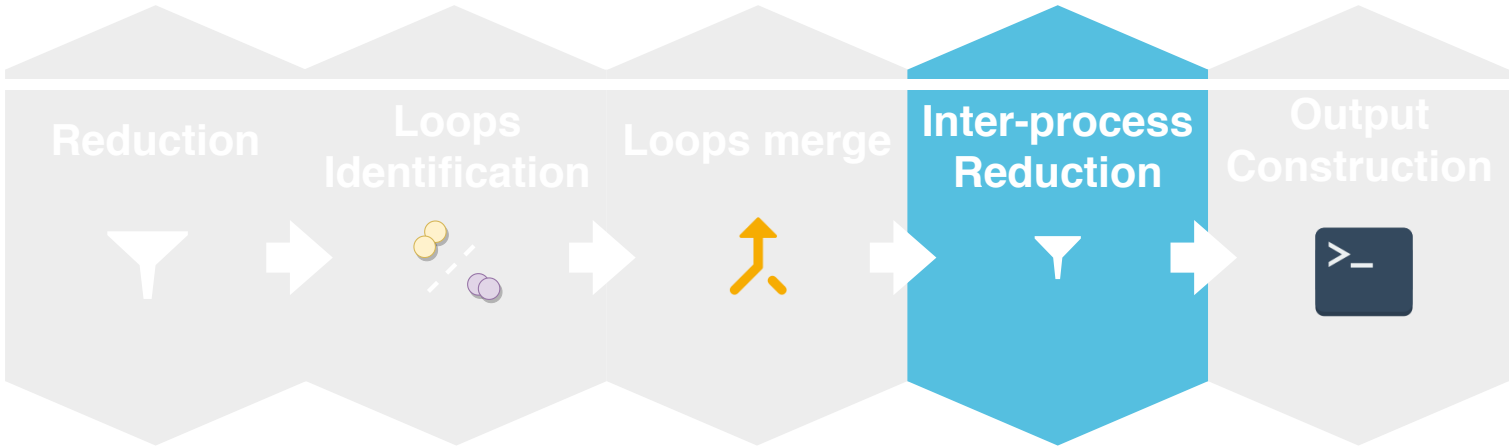

- Colapse same MPI call site from different processes

Workflow

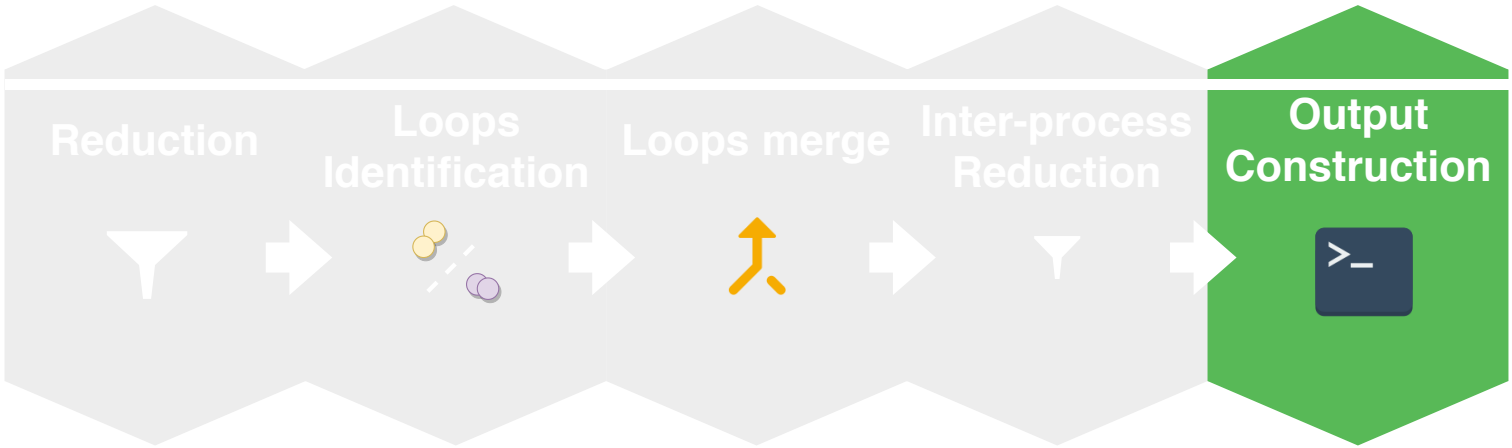

- Construction of the output

Workflow

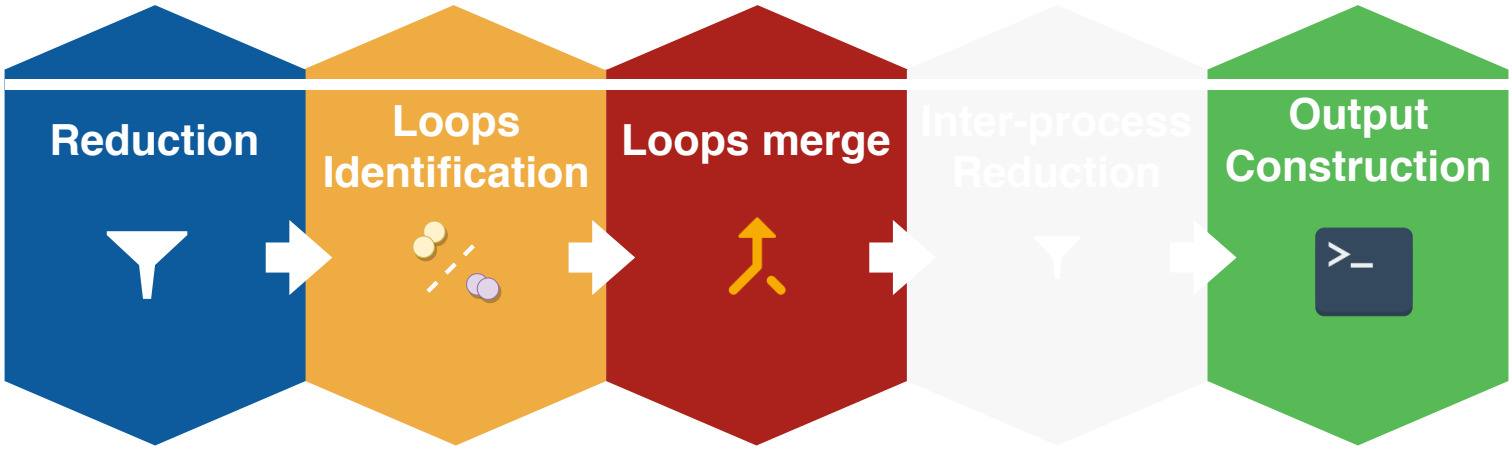

- Parsing and reducing information from trace
- Identify loops by MPI call sites classification
- Find out the hierarchical relations between loops

- Construction of the output

Workflow

Trace reduction step

Input Tracefile, i.e. Sequence of timestamped events ordered by time.
Output Set of unique MPI calls sites with attached information.

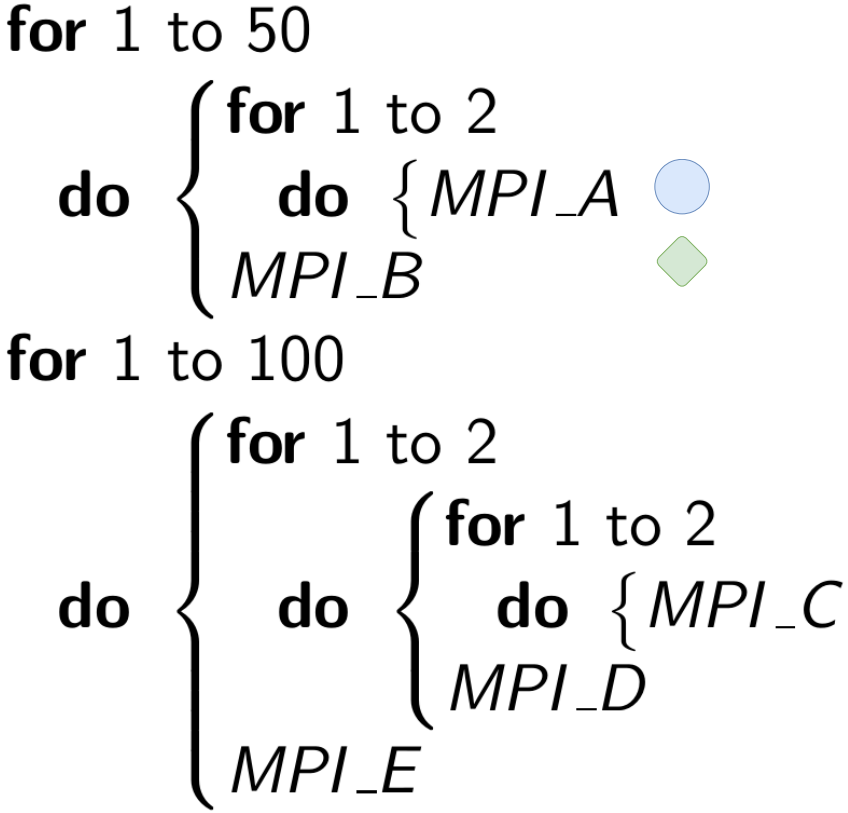

Workflow

Trace reduction step

Input Tracefile, i.e. Sequence of timestamped events ordered by time.
Output Set of unique MPI calls sites with attached information.

for 1 to 50
  do { **for** 1 to 2
    do { *MPI_A*
    MPI_B

for 1 to 100
  do { **for** 1 to 2
    do { **for** 1 to 2
      do { *MPI_C*
      MPI_D
    MPI_E

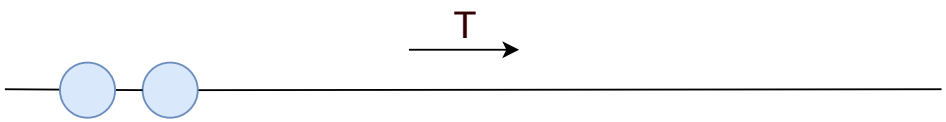

Workflow

Trace reduction step

Input Tracefile, i.e. Sequence of timestamped events ordered by time.
Output Set of unique MPI calls sites with attached information.

for 1 to 50
  do { **for** 1 to 2
    do { *MPI_A*
    MPI_B

for 1 to 100
  do { **for** 1 to 2
    do { **for** 1 to 2
      do { *MPI_C*
      MPI_D
    MPI_E

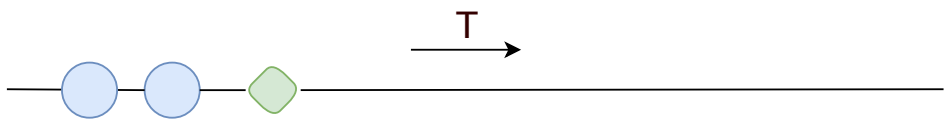

Workflow

Trace reduction step

Input Tracefile, i.e. Sequence of timestamped events ordered by time.
Output Set of unique MPI calls sites with attached information.

for 1 to 50
  do { **for** 1 to 2
    do { *MPI_A*
  MPI_B

for 1 to 100
  do { **for** 1 to 2
    do { **for** 1 to 2
      do { *MPI_C*
    MPI_D
  MPI_E

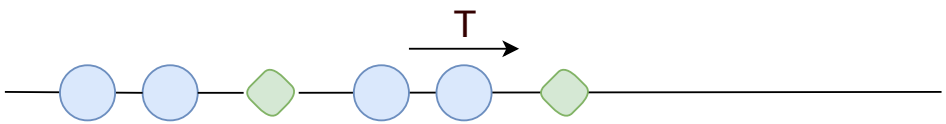

Workflow

Trace reduction step

Input Tracefile, i.e. Sequence of timestamped events ordered by time.
Output Set of unique MPI calls sites with attached information.

for 1 to 50
  do
    for 1 to 2
      do *MPI_A*
    MPI_B

for 1 to 100
  do
    for 1 to 2
      do
        for 1 to 2
          do *MPI_C*
        MPI_D
    MPI_E

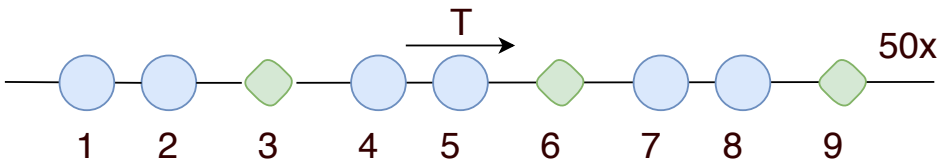

Workflow

Trace reduction step

Input Tracefile, i.e. Sequence of timestamped events ordered by time.
Output Set of unique MPI calls sites with attached information.

for 1 to 50
  do
    for 1 to 2
      do { *MPI_A*
    MPI_B
for 1 to 100
  do
    for 1 to 2
      do
        for 1 to 2
          do { *MPI_C*
        MPI_D
    MPI_E

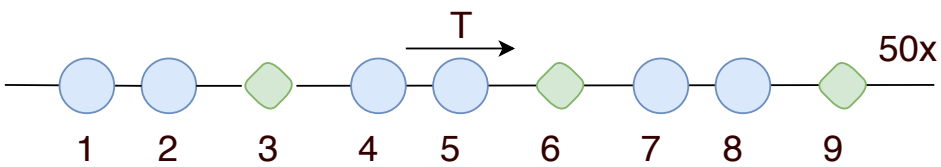

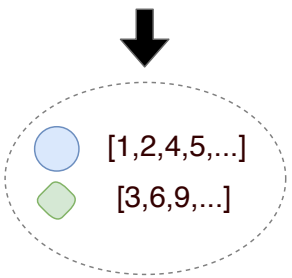

Workflow

Trace reduction step

- Is a sort of Map & Reduce
- Every MPI call is identified by its **signature**: Ordered sequence of pairs (*file*, *line*) that define the call path, i.e. The dynamic position
- Additionally **less representative MPI calls are filtered** (10%).

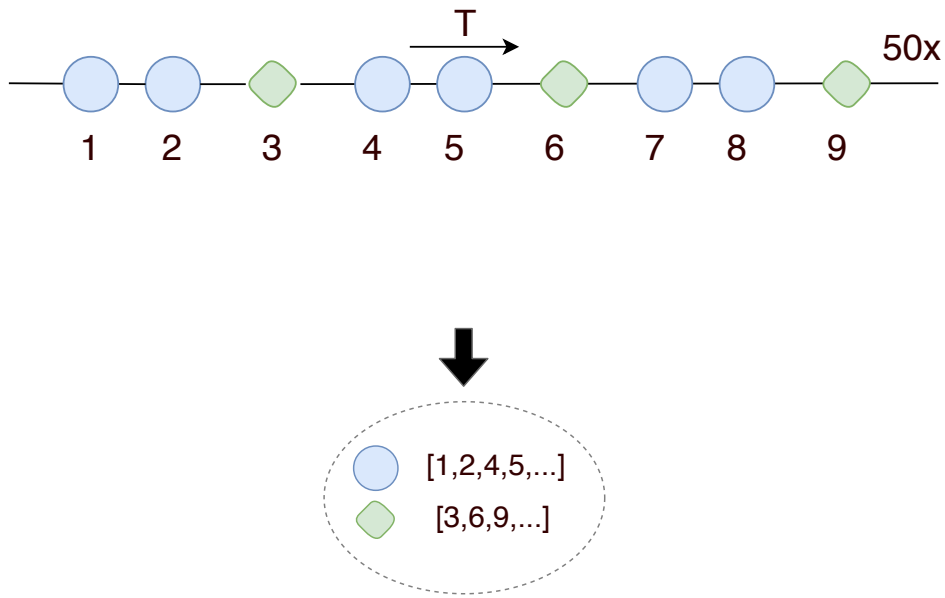

Workflow

Trace reduction step

- Is a sort of Map & Reduce
- Every MPI call is identified by its **signature**: Ordered sequence of pairs (*file*, *line*) that define the call path, i.e. The dynamic position
- Additionally **less representative MPI calls are filtered** (10%).

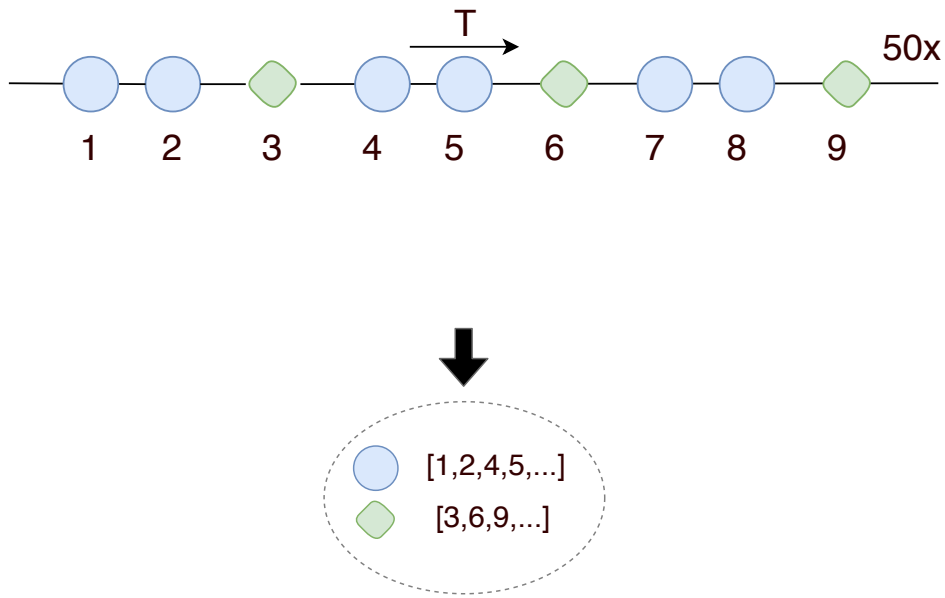

Keynote

Trace reduction is **the key step for scalability** because repetitivity of HPC applications.

Workflow

Loops identification

Input Set of unique MPI calls sites with attached information.
Output Set of clusters of MPI call sites → Loops

for 1 to 50
  do {
    for 1 to 2
      do { *MPI_A*
    MPI_B

for 1 to 100
  do {
    for 1 to 2
      do {
        for 1 to 2
          do { *MPI_C*
        MPI_D
    MPI_E

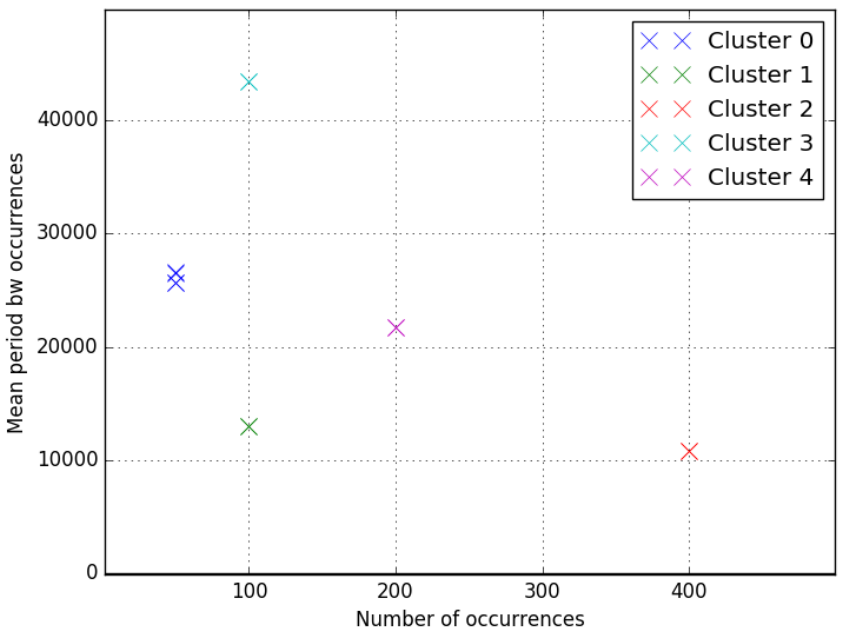

Workflow

Loops identification

Input Set of unique MPI calls sites with attached information.
Output Set of clusters of MPI call sites → Loops

for 1 to 50
  do
    for 1 to 2
      do *MPI_A*
    MPI_B
for 1 to 100
  do
    for 1 to 2
      do
        for 1 to 2
          do *MPI_C*
        MPI_D
    MPI_E

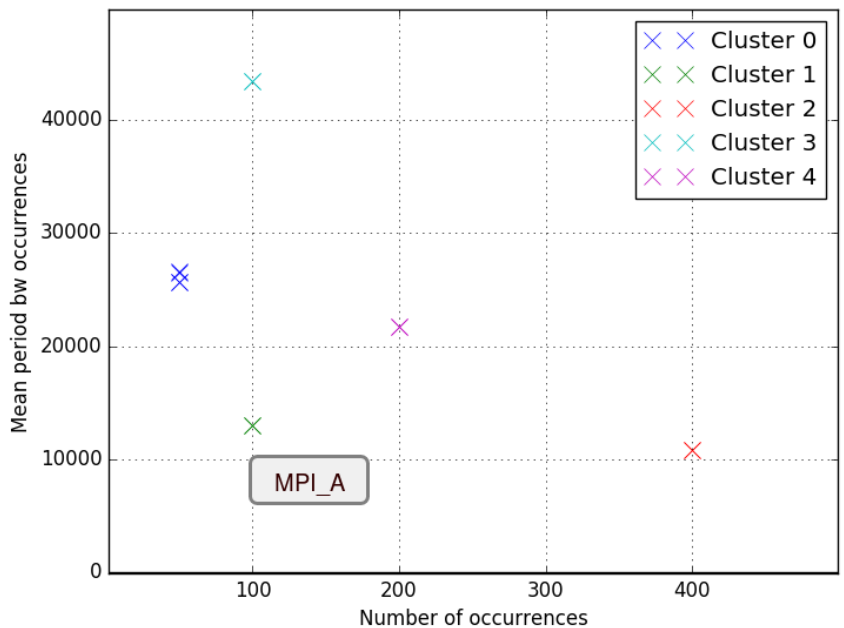

Workflow

Loops identification

Input Set of unique MPI calls sites with attached information.
Output Set of clusters of MPI call sites → Loops

for 1 to 50
  do
    for 1 to 2
      do *MPI_A*
    MPI_B

for 1 to 100
  do
    for 1 to 2
      do
        for 1 to 2
          do *MPI_C*
        MPI_D
    MPI_E

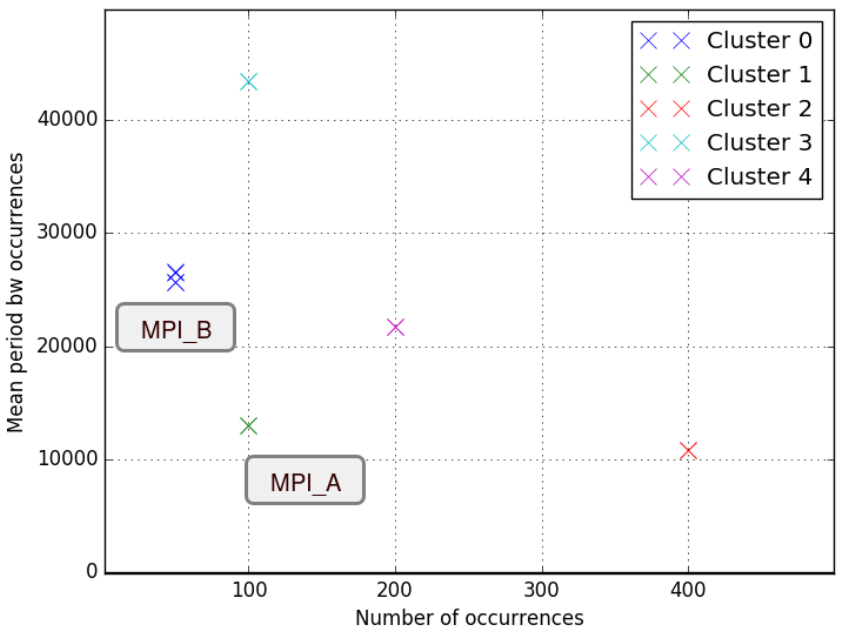

Workflow

Loops identification

Input Set of unique MPI calls sites with attached information.
Output Set of clusters of MPI call sites → Loops

for 1 to 50
  do {
    for 1 to 2
      do { *MPI_A*
    MPI_B
for 1 to 100
  do {
    for 1 to 2
      do {
        for 1 to 2
          do { *MPI_C*
        MPI_D
    MPI_E

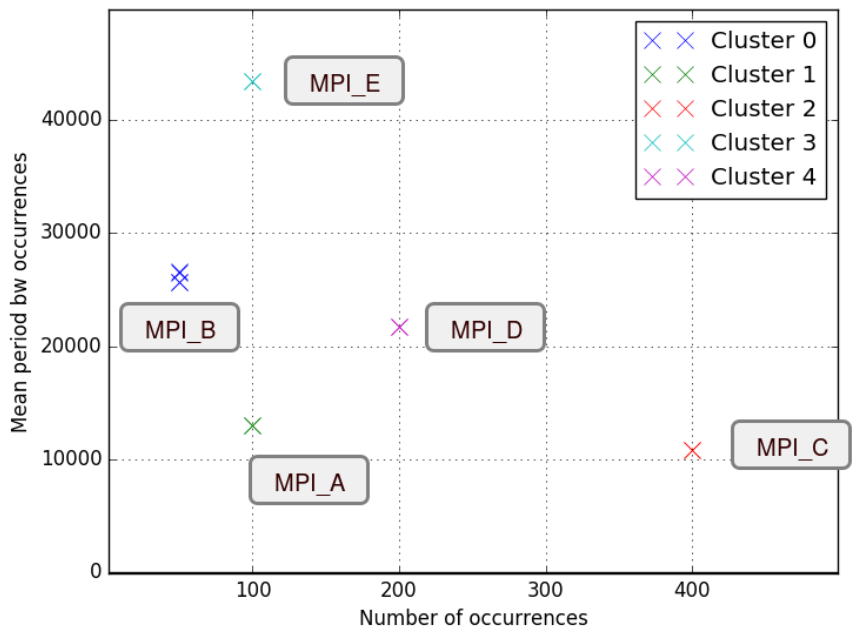

Workflow

Loops identification

- Need for some classification algorithm
- **DBSCAN** as clustering algorithm (preferred to K-means)
  - ϵ empirically set to 0.2 (in general)
  - minPts set to 1
- Resulting clusters **will be considered loops**.
  - Number of repetitions $\rightarrow$ Number of iterations.
  - Mean time between repetitions $\rightarrow$ Mean iterations time.

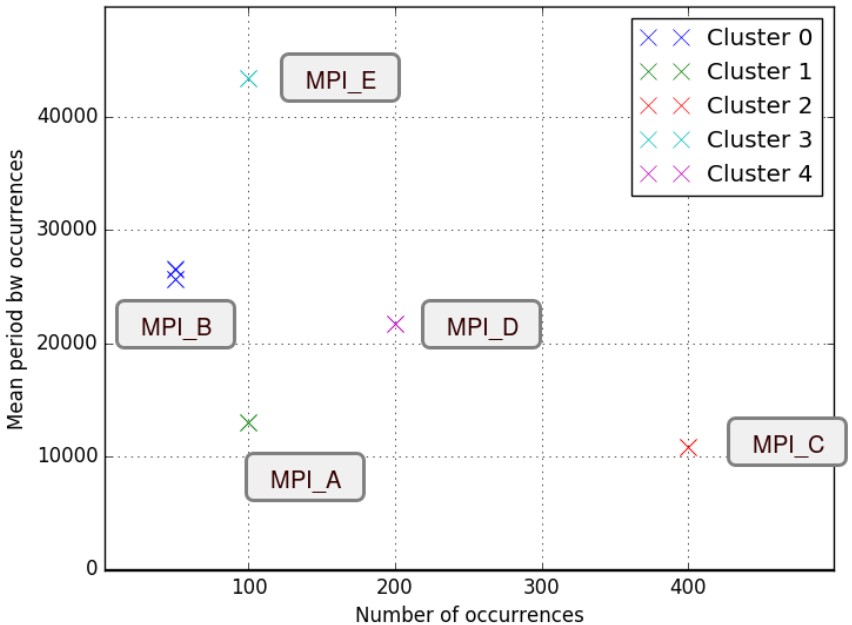

Workflow

Loops merge

Input Set of loops.
Output Set of top level loops with its related nested loops.

Workflow

Loops merge

Input Set of loops.
Output Set of top level loops with its related nested loops.

Intuition

- Isolated loops are just **pieces of the overall puzzle**.
- By discovering its hierarchical relations **the structure of the application will be revealed**...
- As well as the different phases

Workflow

Loops merge

Input Set of loops.
Output Set of top level loops with its related nested loops.

Intuition

- Isolated loops are just **pieces of the overall puzzle**.
- By discovering its hierarchical relations **the structure of the application will be revealed**...
- As well as the different phases

We can assume

- Outer loop will have **less iterations** than nested one.
- Outer loop will spend **more time** per iteration.

Workflow

Loops merge

for 1 to 50

do { **for** 1 to 2 **do** { *MPI_A* ; *MPI_B* }

for 1 to 100

do { **for** 1 to 2 **do** { **for** 1 to 2 **do** { *MPI_C* ; *MPI_D* } } ; *MPI_E* }

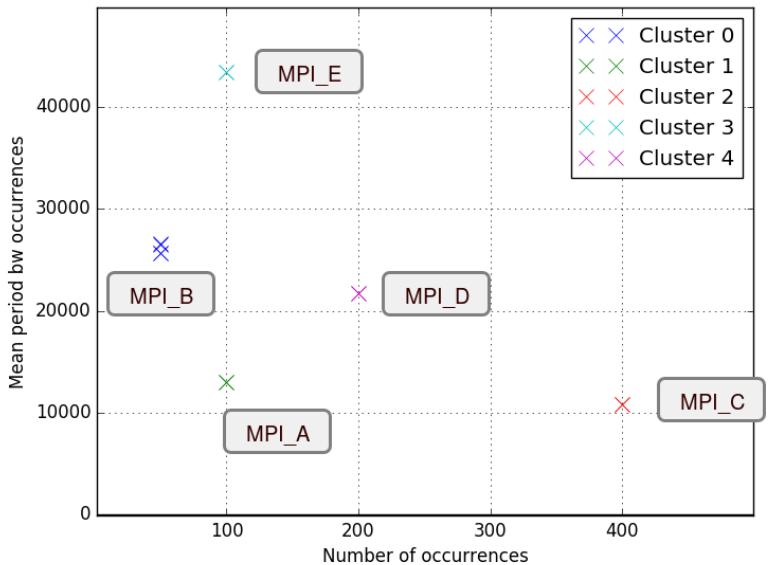

Workflow

Loops merge

for 1 to 50
  do { **for** 1 to 2
    do { *MPI_A*
    MPI_B
for 1 to 100
  do { **for** 1 to 2
    do { **for** 1 to 2
      do { *MPI_C*
      MPI_D
    MPI_E

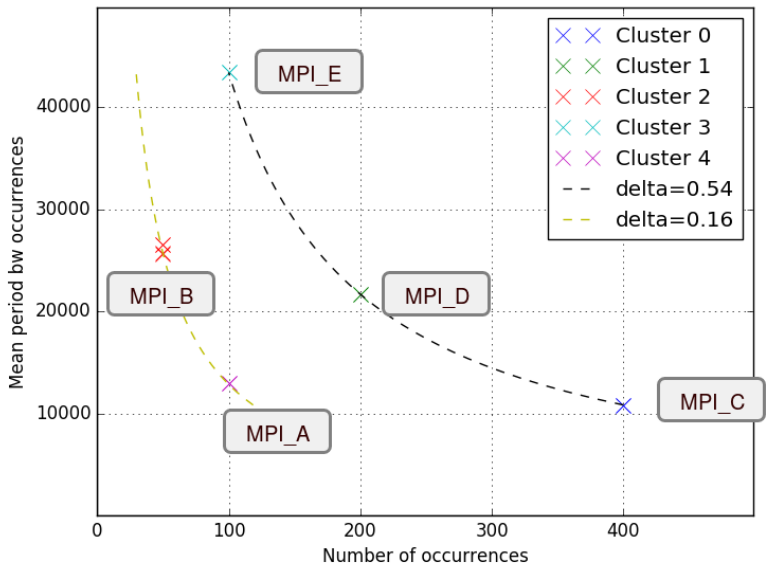

- Classify loops by its delta
  - All lies on same $f(x) = \frac{\delta * T_{exe}}{x}$ being $0 < \delta < 1$
- And then merge them in a less to more iterations fashion
  - Always checking iterations interleaving

Workflow

Loops merge

for 1 to 50
  do { **for** 1 to 2
    do { *MPI_A*
    MPI_B
for 1 to 100
  do { **for** 1 to 2
    do { **for** 1 to 2
      do { *MPI_C*
      MPI_D
    MPI_E

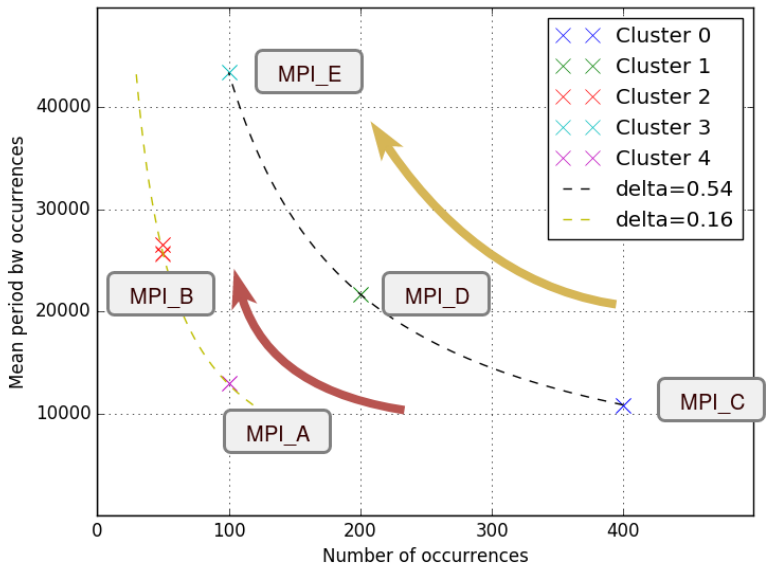

- Classify loops by its delta
  - All lies on same $f(x) = \frac{\delta * T_{exe}}{x}$ being $0 < \delta < 1$
- And then merge them in a less to more iterations fashion
  - Always checking iterations interleaving

Workflow

Loops merge

for 1 to 50
  do { **for** 1 to 2
    do { *MPI_A*
    MPI_B
for 1 to 100
  do { **for** 1 to 2
    do { **for** 1 to 2
      do { *MPI_C*
      MPI_D
    MPI_E

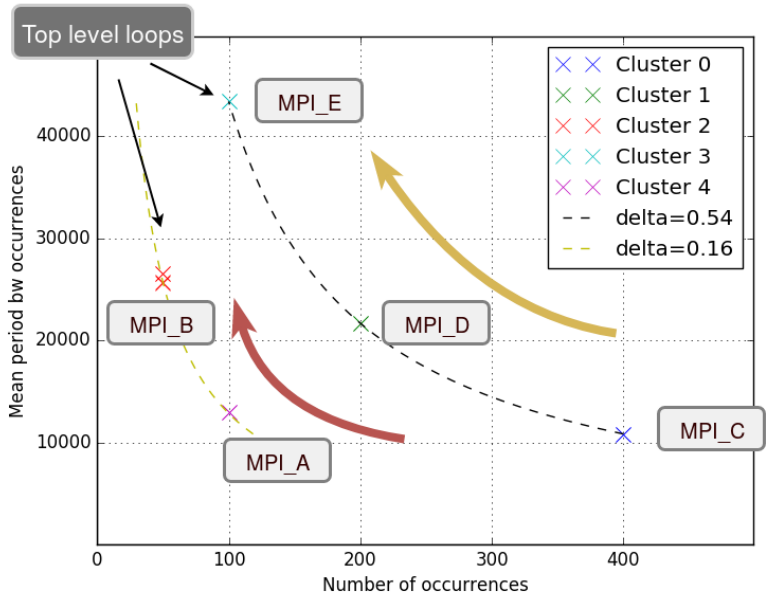

- Classify loops by its delta
  - All lies on same $f(x) = \frac{\delta * T_{exe}}{x}$ being $0 < \delta < 1$
- And then merge them in a less to more iterations fashion
  - Always checking iterations interleaving

Keynote

This mechanism is used for phases detection.

Workflow

Pseudocode construction

Input Set of top level loops with rank conditional structures.
Output Pseudocode representing the actual application structure.

Workflow

Pseudocode construction

Input Set of top level loops with rank conditional structures.
Output Pseudocode representing the actual application structure.

- Straightforward construction
- Just a refination of the data is needed
  1. Extracting common call path levels from code block (loops and conditional blocks)
  2. Removing **contiguous repetitive information** what has not been removed in previous step

Workflow

Pseudocode construction

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 50	-	-	-
test-13.c	17	: : MPI_Barrier(CommId:1)	26.43us	-	0.3
		: : FOR 1 TO 2.0	-	-	-
		: : : IF rank in [1]	-	-	-
test-13.c	23	: : : : MPI_Send(1:0)	22.8us	4.0B	0.53
		: : : IF rank in [0]	-	-	-
test-13.c	25	: : : : MPI_Recv()	27.09us	4.0B	0.54
		: : END LOOP	-	-	-
test-13.c	**28**	**: : [~ computation ~]**	**10.18ms**	**-**	**0.06**
test-13.c	28	: : MPI_Barrier(CommId:1)	106.13us	-	1.43
		: END LOOP	-	-	-
	0	main()	-	-	-
		: FOR 1 TO 100	-	-	-
		: : FOR 1 TO 2.0	-	-	-
		: : : FOR 1 TO 2.0	-	-	-
		: : : : IF rank in [1]	-	-	-
test-13.c	40	: : : : : MPI_Send(1:0)	23.0us	4.0B	0.5
		: : : : IF rank in [0]	-	-	-
test-13.c	42	: : : : : MPI_Recv()	25.1us	4.0B	0.51
		: : : END LOOP	-	-	-
test-13.c	45	: : : MPI_Barrier(CommId:1)	23.93us	-	0.84
		: : END LOOP	-	-	-
test-13.c	**48**	**: : [~ computation ~]**	**50.2ms**	**-**	**0.07**
test-13.c	48	: : MPI_Barrier(CommId:1)	124.99us	-	1.62
		: END LOOP	-	-	-

Workflow

Pseudocode construction

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 50	-	-	-
test-13.c	17	: : MPI_Barrier(CommId:1)	26.43us	-	0.3
		: : FOR 1 TO 2.0	-	-	-
		: : : IF rank in [1]	-	-	-
test-13.c	23	: : : : MPI_Send(1:0)	22.8us	4.0B	0.53
		: : : IF rank in [0]	-	-	-
test-13.c	25	: : : : MPI_Recv()	27.09us	4.0B	0.54
		: : END LOOP	-	-	-
test-13.c	**28**	**: : [~ computation ~]**	**10.18ms**	**-**	**0.06**
test-13.c	28	: : MPI_Barrier(CommId:1)	106.13us	-	1.43
		: END LOOP	-	-	-
	0	main()	-	-	-
		: FOR 1 TO 100	-	-	-
		: : FOR 1 TO 2.0	-	-	-
		: : : FOR 1 TO 2.0	-	-	-
		: : : : IF rank in [1]	-	-	-
test-13.c	40	: : : : : MPI_Send(1:0)	23.0us	4.0B	0.5
		: : : : IF rank in [0]	-	-	-
test-13.c	42	: : : : : MPI_Recv()	25.1us	4.0B	0.51
		: : : END LOOP	-	-	-
test-13.c	45	: : : MPI_Barrier(CommId:1)	23.93us	-	0.84
		: : END LOOP	-	-	-
test-13.c	**48**	**: : [~ computation ~]**	**50.2ms**	**-**	**0.07**
test-13.c	48	: : MPI_Barrier(CommId:1)	124.99us	-	1.62
		: END LOOP	-	-	-

Workflow

Pseudocode construction

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 50	-	-	-
test-13.c	17	: : MPI_Barrier(CommId:1)	26.43us	-	0.3
		: : FOR 1 TO 2.0	-	-	-
		: : : IF rank in [1]	-	-	-
test-13.c	23	: : : : MPI_Send(1:0)	22.8us	4.0B	0.53
		: : : IF rank in [0]	-	-	-
test-13.c	25	: : : : MPI_Recv()	27.09us	4.0B	0.54
		: : END LOOP	-	-	-
test-13.c	**28**	**: : [~ computation ~]**	**10.18ms**	**-**	**0.06**
test-13.c	28	: : MPI_Barrier(CommId:1)	106.13us	-	1.43
		: END LOOP	-	-	-
	0	main()	-	-	-
		: FOR 1 TO 100	-	-	-
		: : FOR 1 TO 2.0	-	-	-
		: : : FOR 1 TO 2.0	-	-	-
		: : : : IF rank in [1]	-	-	-
test-13.c	40	: : : : : MPI_Send(1:0)	23.0us	4.0B	0.5
		: : : : IF rank in [0]	-	-	-
test-13.c	42	: : : : : MPI_Recv()	25.1us	4.0B	0.51
		: : : END LOOP	-	-	-
test-13.c	45	: : : MPI_Barrier(CommId:1)	23.93us	-	0.84
		: : END LOOP	-	-	-
test-13.c	**48**	**: : [~ computation ~]**	**50.2ms**	**-**	**0.07**
test-13.c	48	: : MPI_Barrier(CommId:1)	124.99us	-	1.62
		: END LOOP	-	-	-

Workflow

Pseudocode construction

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 50	-	-	-
test-13.c	17	: : MPI_Barrier(CommId:1)	26.43us	-	0.3
		: : FOR 1 TO 2.0	-	-	-
		: : : IF rank in [1]	-	-	-
test-13.c	23	: : : : MPI_Send(1:0)	22.8us	4.0B	0.53
		: : : IF rank in [0]	-	-	-
test-13.c	25	: : : : MPI_Recv()	27.09us	4.0B	0.54
		: : END LOOP	-	-	-
test-13.c	**28**	**: : [~ computation ~]**	**10.18ms**	**-**	**0.06**
test-13.c	28	: : MPI_Barrier(CommId:1)	106.13us	-	1.43
		: END LOOP	-	-	-
	0	main()	-	-	-
		: FOR 1 TO 100	-	-	-
		: : FOR 1 TO 2.0	-	-	-
		: : : FOR 1 TO 2.0	-	-	-
		: : : : IF rank in [1]	-	-	-
test-13.c	40	: : : : : MPI_Send(1:0)	23.0us	4.0B	0.5
		: : : : IF rank in [0]	-	-	-
test-13.c	42	: : : : : MPI_Recv()	25.1us	4.0B	0.51
		: : : END LOOP	-	-	-
test-13.c	45	: : : MPI_Barrier(CommId:1)	23.93us	-	0.84
		: : END LOOP	-	-	-
test-13.c	**48**	**: : [~ computation ~]**	**50.2ms**	**-**	**0.07**
test-13.c	48	: : MPI_Barrier(CommId:1)	124.99us	-	1.62
		: END LOOP	-	-	-

Workflow

Pseudocode construction

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 50	-	-	-
test-13.c	17	: : MPI_Barrier(CommId:1)	26.43us	-	0.3
		: : FOR 1 TO 2.0	-	-	-
		: : : IF rank in [1]	-	-	-
test-13.c	23	: : : : MPI_Send(1:0)	22.8us	4.0B	0.53
		: : : IF rank in [0]	-	-	-
test-13.c	25	: : : : MPI_Recv()	27.09us	4.0B	0.54
		: : END LOOP	-	-	-
test-13.c	**28**	**: : [~ computation ~]**	**10.18ms**	**-**	**0.06**
test-13.c	28	: : MPI_Barrier(CommId:1)	106.13us	-	1.43
		: END LOOP	-	-	-
	0	main()	-	-	-
		: FOR 1 TO 100	-	-	-
		: : FOR 1 TO 2.0	-	-	-
		: : : FOR 1 TO 2.0	-	-	-
		: : : : IF rank in [1]	-	-	-
test-13.c	40	: : : : : MPI_Send(1:0)	23.0us	4.0B	0.5
		: : : : IF rank in [0]	-	-	-
test-13.c	42	: : : : : MPI_Recv()	25.1us	4.0B	0.51
		: : : END LOOP	-	-	-
test-13.c	45	: : : MPI_Barrier(CommId:1)	23.93us	-	0.84
		: : END LOOP	-	-	-
test-13.c	**48**	**: : [~ computation ~]**	**50.2ms**	**-**	**0.07**
test-13.c	48	: : MPI_Barrier(CommId:1)	124.99us	-	1.62
		: END LOOP	-	-	-

```
(struct_detection)> help

Documented commands (type help <topic>):
========================================
clustering  help  paraver  pseudocode  q  quit

(struct_detection)> help pseudocode
pseudocode commands:
wo-cs: Do not show call paths
w-burst-info: Show CPU burst information
w-burst-threshold: Show burst above thresshold
default: Default information
ranks: Show pseudocode just for rank
(struct_detection)>
```

Refination

Clustering aliasing $\rightarrow$ Same behavior

for $1 to 100$

do $\begin{cases} MPI\_A() \\ MPI\_B() \end{cases}$

for $1 to 100$

do $\begin{cases} MPI\_C() \\ MPI\_D() \end{cases}$

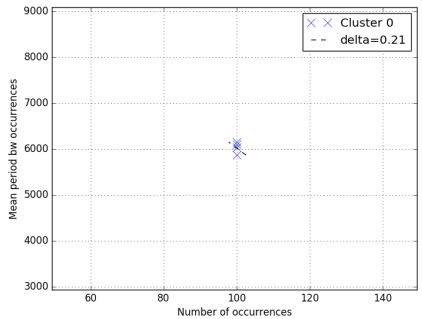

Refination

Clustering aliasing $\rightarrow$ Same behavior

$$\textbf{for } 1 to 100 \textbf{ do } \begin{cases} MPI\_A() \\ MPI\_B() \end{cases}$$
$$\textbf{for } 1 to 100 \textbf{ do } \begin{cases} MPI\_C() \\ MPI\_D() \end{cases}$$

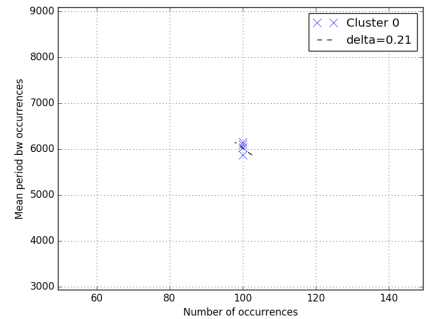

Clustering split $\rightarrow$ Data conditions

$$\textbf{for } i = 1 \text{ to } 10 \textbf{ do } \begin{cases} MPI\_B() \\ \textbf{if } \quad odd(i) \\ \quad \textbf{then } MPI\_A() \\ MPI\_B() \end{cases}$$

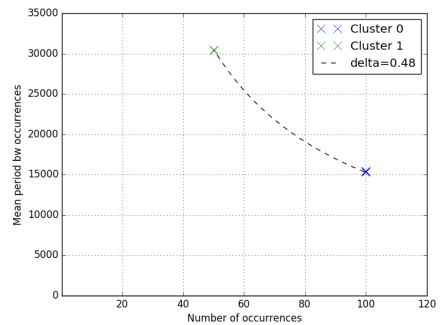

Refination

Clustering aliasing $\rightarrow$ Same behavior

$$\textbf{for } 1 to 100 \; \textbf{do} \begin{cases} MPI\_A() \\ MPI\_B() \end{cases}$$
$$\textbf{for } 1 to 100 \; \textbf{do} \begin{cases} MPI\_C() \\ MPI\_D() \end{cases}$$

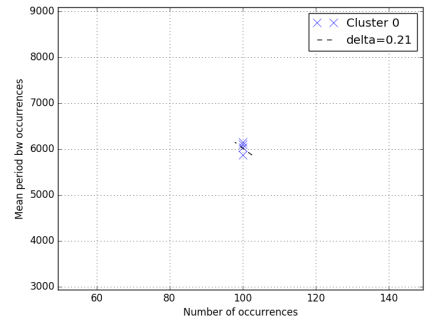

Clustering split $\rightarrow$ Data conditions

$$\textbf{for } i = 1 \text{ to } 10 \; \textbf{do} \begin{cases} MPI\_B() \\ \textbf{if} \quad odd(i) \\ \quad \textbf{then } MPI\_A() \\ MPI\_B() \end{cases}$$

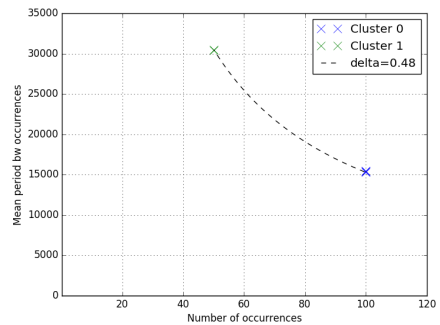

Keynote

Both **solved** but better to **avoid them**

Outline

Best features analysis

Are the used features **the best**?

- Number of repetitions and Mean time between repetitions

[1] Strictly limited by available hardware

Best features analysis

Are the used features **the best**?

- Number of repetitions and Mean time between repetitions

Can other **additional/alternative** features prevent clustering split/aliasing?

- IPC, %Loads, %(Cond, Uncond, Total) branches [1]

[1] Strictly limited by available hardware

Best features analysis

Are the used features **the best**?

- Number of repetitions and Mean time between repetitions

Can other **additional/alternative** features prevent clustering split/aliasing?

- IPC, %Loads, %(Cond, Uncond, Total) branches [1]

Used techniques:

- Principal Component Analysis
- Random Forest Variable Importance $\rightarrow$ Needs labeled MPI call sites with loop identification.

[1] Strictly limited by available hardware

Best features analysis

Are the used features **the best**?

- Number of repetitions and Mean time between repetitions

Can other **additional/alternative** features prevent clustering split/aliasing?

- IPC, %Loads, %(Cond, Uncond, Total) branches [1]

Used techniques:

- Principal Component Analysis
- Random Forest Variable Importance $\rightarrow$ Needs labeled MPI call sites with loop identification.

Data acquisition:

- Used tracing library (Extrae) is monitorizing calls to shared libraries
- Loops monitors injection by **Mercurium**

[1] Strictly limited by available hardware

Best features analysis

Analysis of results: **Variable Importance**

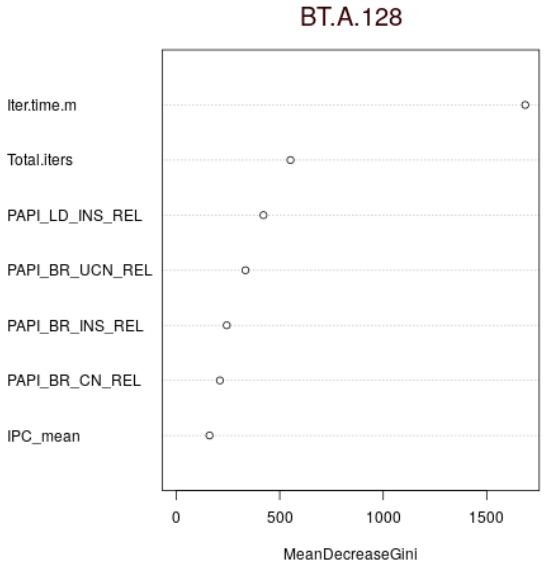

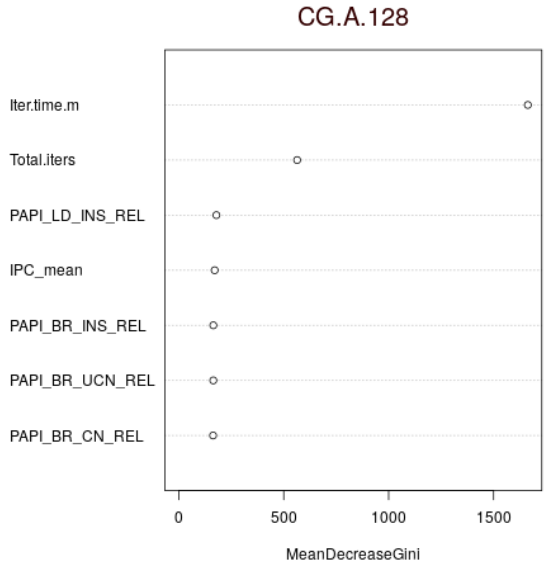

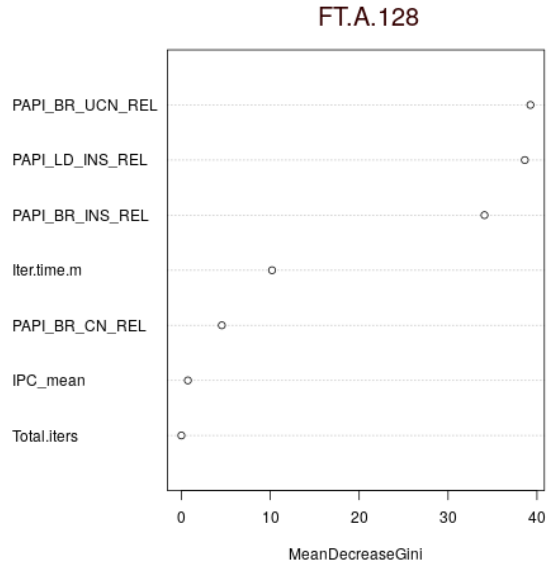

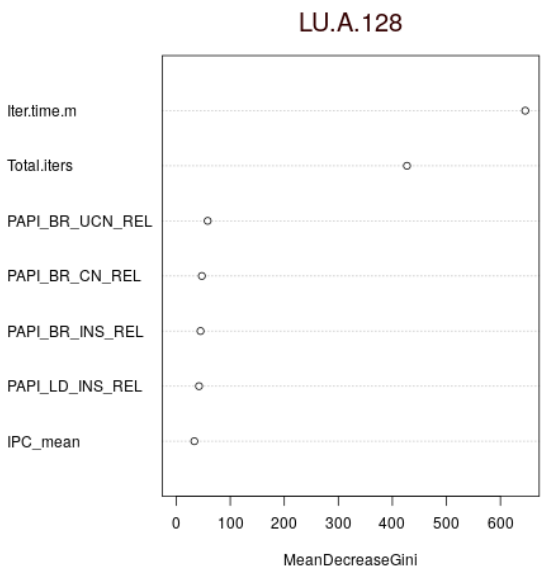

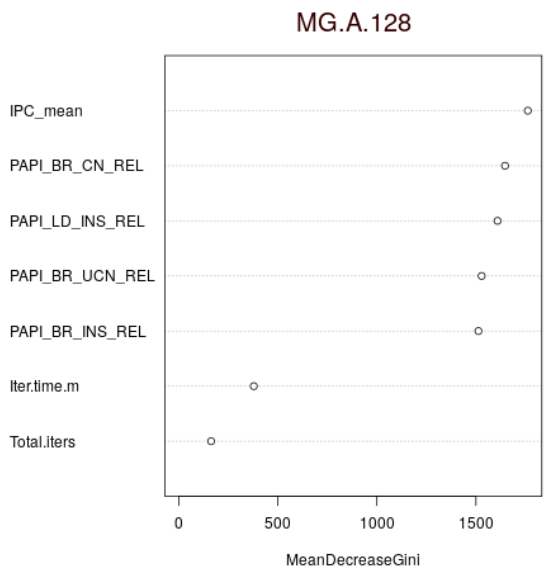

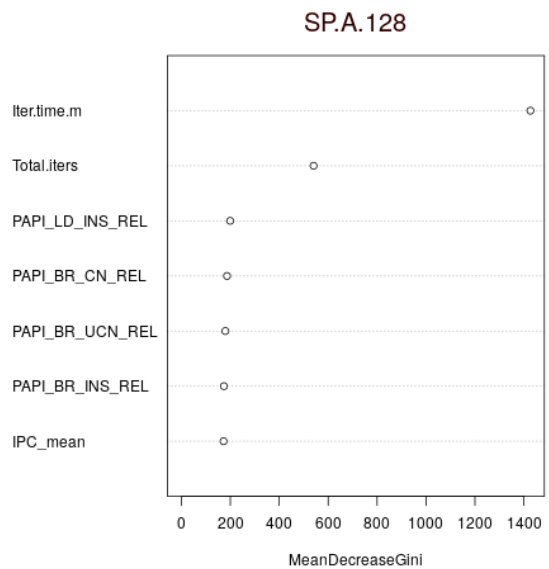

Best features analysis

Analysis of results: **Variable Importance**

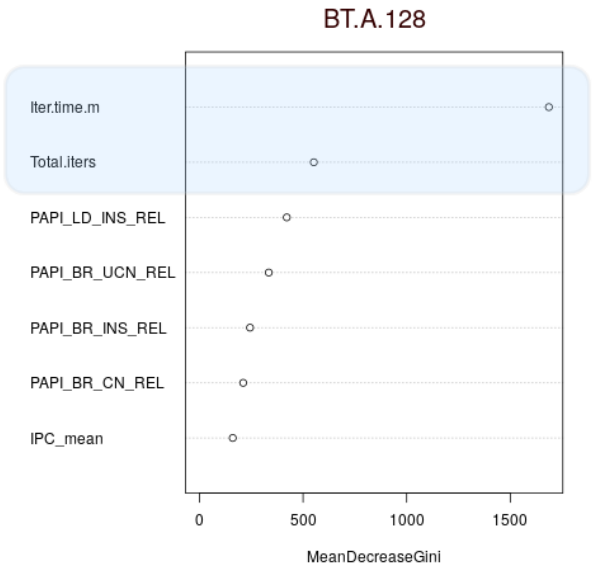

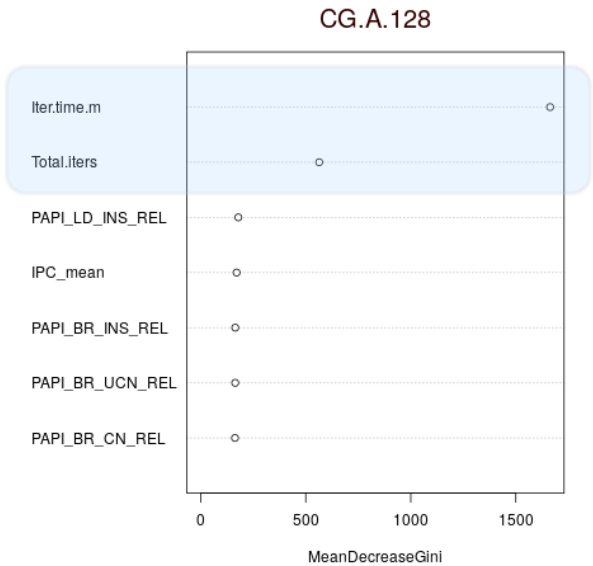

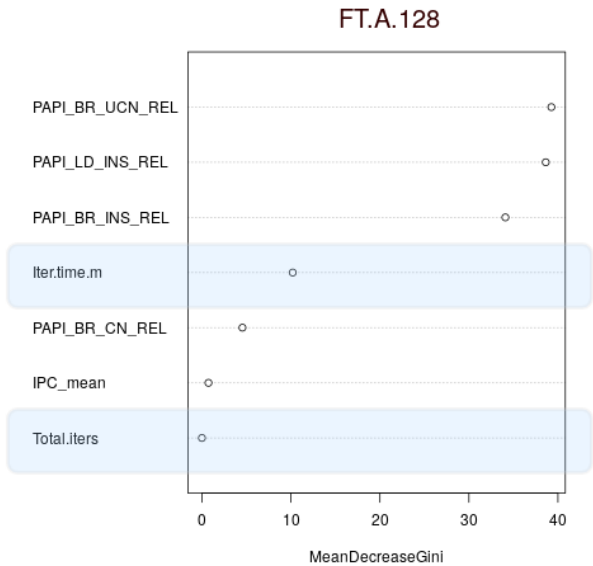

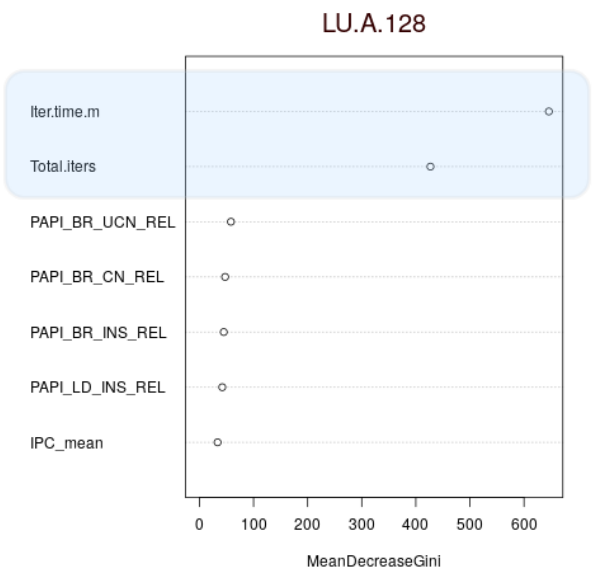

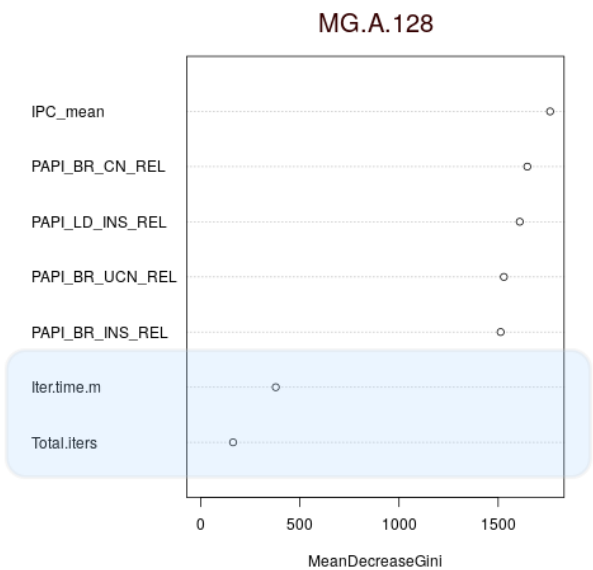

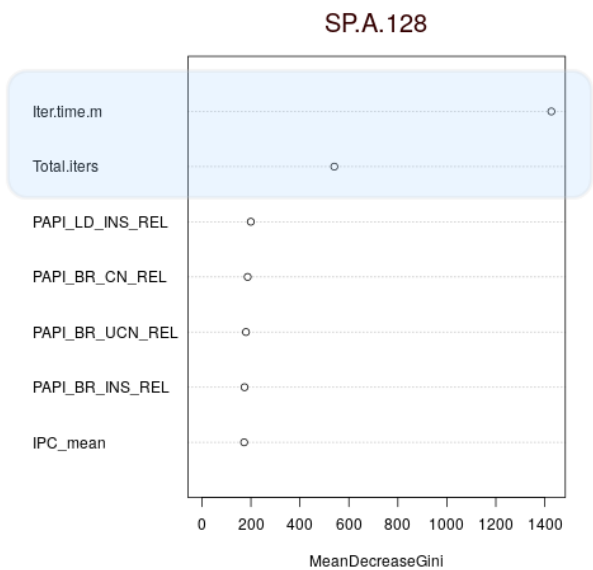

Outline

Validation

Lulesh 2.0

Execution:

- 6014 code lines (C++) (142 MPI calls)
- 125 parallel processes
- $\approx 160MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 6ms$

Validation

Lulesh 2.0

Execution:

- 6014 code lines (C++) (142 MPI calls)
- 125 parallel processes
- $\approx 160MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 6ms$

Results:

- 30 iterations loop

```
main()
: FOR 1 TO 30 [id=0.0]
: : TimeIncrement()
```

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
lulesh.cc	1263	: : : : CalcForceForNodes()	-	-	-
lulesh.cc	1137	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh-comm.cc	436	: : : : : MPI_Isend(0:25)	14.95us	22.52KB	1.9
lulesh-comm.cc	477	: : : : : MPI_Isend(0:5)	10.07us	22.52KB	1.65
lulesh-comm.cc	518	: : : : : MPI_Isend(0:1)	10.5us	22.52KB	0.93
lulesh-comm.cc	589	: : : : : MPI_Isend(0:6)	14.79us	744.0B	1.68
lulesh-comm.cc	606	: : : : : MPI_Isend(0:30)	12.45us	744.0B	1.86
lulesh-comm.cc	623	: : : : : MPI_Isend(0:26)	13.13us	744.0B	1.66
lulesh-comm.cc	837	: : : : : MPI_Isend(0:31)	9.48us	24.0B	0.84
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	541.3us	-	0.43
lulesh.cc	1161	: : : : CommSBN()	-	-	-
lulesh-comm.cc	889	: : : : : MPI_Comm_rank(0:)	9.32us	-	0.49
lulesh-comm.cc	911	: : : : : MPI_Wait(0:)	9.25us	-	0.6
lulesh-comm.cc	945	: : : : : MPI_Wait(0:)	8.6us	-	2.48
lulesh-comm.cc	980	: : : : : MPI_Wait(0:)	8.52us	-	2.35
lulesh-comm.cc	1039	: : : : : MPI_Wait(0:)	8.41us	-	1.72
lulesh-comm.cc	1053	: : : : : MPI_Wait(0:)	8.19us	-	1.64
lulesh-comm.cc	1067	: : : : : MPI_Wait(0:)	8.31us	-	1.98
lulesh-comm.cc	1251	: : : : : MPI_Wait(0:)	58.99us	-	1.44
lulesh.cc	1267	: : : : CommRecv()	-	-	-
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	7.77us	-	0.73
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	9.03us	45.05KB	0.79
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	7.37us	45.05KB	1.02
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	7.34us	45.05KB	0.95
lulesh-comm.cc	192	: : : : : MPI_Irecv(0:6)	7.6us	1.45KB	1.12
lulesh-comm.cc	201	: : : : : MPI_Irecv(0:30)	7.51us	1.45KB	1.33
lulesh-comm.cc	210	: : : : : MPI_Irecv(0:26)	7.42us	1.45KB	1.13
lulesh-comm.cc	345	: : : : : MPI_Irecv(0:31)	7.5us	48.0B	0.95
lulesh.cc	1289	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	10.26us	-	0.35
lulesh.cc	1292	: : : : CommSyncPosVel()	-	-	-
lulesh-comm.cc	1310	: : : : : MPI_Comm_rank(0:)	22.14us	-	0.73
lulesh-comm.cc	1332	: : : : : MPI_Wait(0:)	274.73us	-	0.79
lulesh-comm.cc	1366	: : : : : MPI_Wait(0:)	72.95us	-	2.46
lulesh-comm.cc	1402	: : : : : MPI_Wait(0:)	23.24us	-	2.3
lulesh-comm.cc	1461	: : : : : MPI_Wait(0:)	198.36us	-	1.79
lulesh-comm.cc	1475	: : : : : MPI_Wait(0:)	7.83us	-	1.83
lulesh-comm.cc	1489	: : : : : MPI_Wait(0:)	7.75us	-	2.08
lulesh-comm.cc	1674	: : : : : MPI_Wait(0:)	333.07us	-	1.73
lulesh.cc	2656	: : : LagrangeElements()	-	-	-
lulesh.cc	2461	: : : CalcQForElems()	-	-	-
lulesh.cc	1992	: : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	12.61us	21.09KB	0.43
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	9.23us	21.09KB	1.01
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	8.79us	21.09KB	1.17
lulesh.cc	2010	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	15.24us	-	1.52
lulesh-comm.cc	436	: : : : : MPI_Isend(0:25)	13.14us	21.09KB	1.9
lulesh-comm.cc	477	: : : : : MPI_Isend(0:5)	10.28us	21.09KB	1.89
lulesh-comm.cc	518	: : : : : MPI_Isend(0:1)	10.57us	21.09KB	1.26
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	291.68us	-	0.45
lulesh.cc	2014	: : : : CommMonoQ()	-	-	-
lulesh-comm.cc	1734	: : : : : MPI_Comm_rank(0:)	9.51us	-	0.45
lulesh-comm.cc	1757	: : : : : MPI_Wait(0:)	8.88us	-	0.81
lulesh-comm.cc	1791	: : : : : MPI_Wait(0:)	8.78us	-	2.04
lulesh-comm.cc	1824	: : : : : MPI_Wait(0:)	8.53us	-	2.13
		: END LOOP	-	-	-

Validation

Lulesh 2.0

Execution:

- 6014 code lines (C++) (142 MPI calls)
- 125 parallel processes
- $\approx 160MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 6ms$

Results:

- 30 iterations loop
- 5 long CPU burst

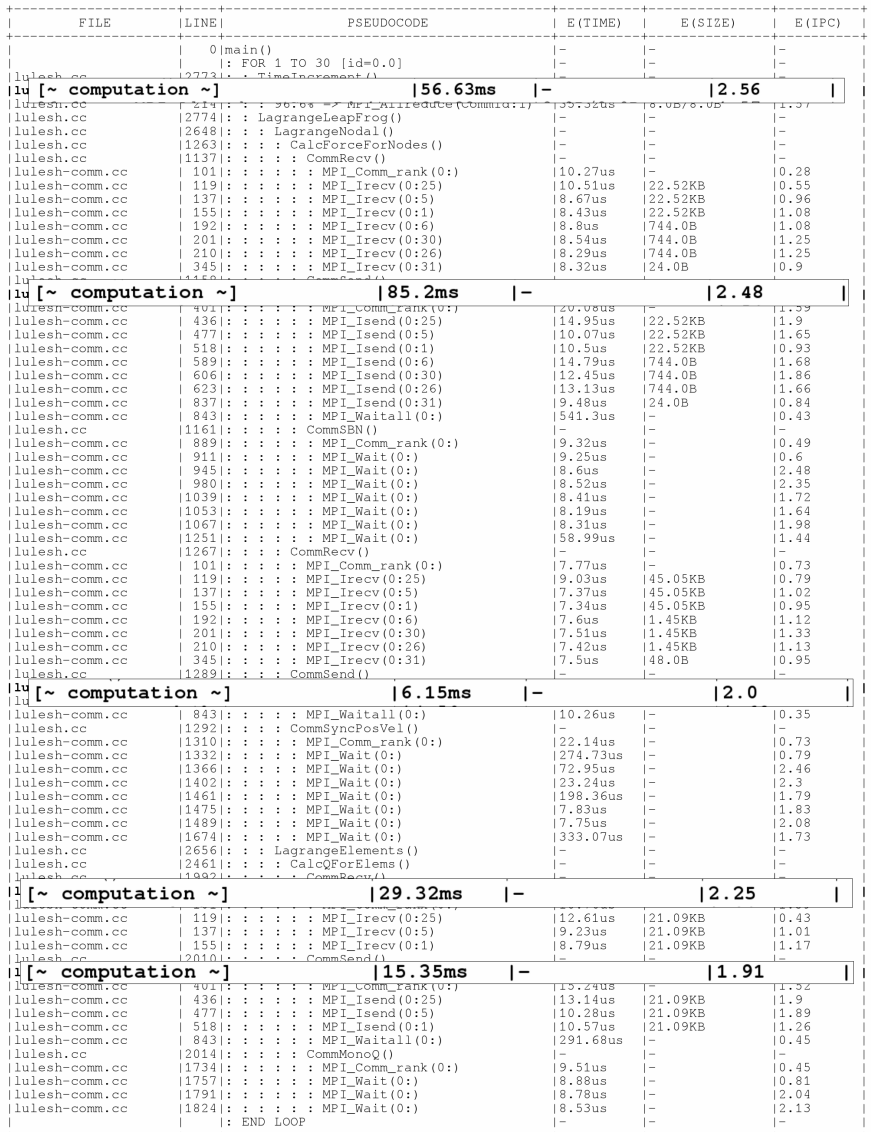

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 30 [id=0.0]	-	-	-
[~ computation ~]		56.63ms	-	2.56	
lulesh.cc	2774	: : LagrangeLeapFrog()	-	-	-
lulesh.cc	2648	: : : LagrangeNodal()	-	-	-
lulesh.cc	1263	: : : : CalcForceForNodes()	-	-	-
lulesh.cc	1137	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	101	: : : : : : MPI_Comm_rank(0:)	10.27us	-	0.28
lulesh-comm.cc	119	: : : : : : MPI_Irecv(0:25)	10.51us	22.52KB	0.55
lulesh-comm.cc	137	: : : : : : MPI_Irecv(0:5)	8.67us	22.52KB	0.96
lulesh-comm.cc	155	: : : : : : MPI_Irecv(0:1)	8.43us	22.52KB	1.08
lulesh-comm.cc	192	: : : : : : MPI_Irecv(0:6)	8.8us	744.0B	1.08
lulesh-comm.cc	201	: : : : : : MPI_Irecv(0:30)	8.54us	744.0B	1.25
lulesh-comm.cc	210	: : : : : : MPI_Irecv(0:26)	8.29us	744.0B	1.25
lulesh-comm.cc	345	: : : : : : MPI_Irecv(0:31)	8.32us	24.0B	0.9
[~ computation ~]		85.2ms	-	2.48	
lulesh-comm.cc	436	: : : : : : MPI_Isend(0:25)	14.95us	22.52KB	1.9
lulesh-comm.cc	477	: : : : : : MPI_Isend(0:5)	10.07us	22.52KB	1.65
lulesh-comm.cc	518	: : : : : : MPI_Isend(0:1)	10.5us	22.52KB	0.93
lulesh-comm.cc	589	: : : : : : MPI_Isend(0:6)	14.79us	744.0B	1.68
lulesh-comm.cc	606	: : : : : : MPI_Isend(0:30)	12.45us	744.0B	1.86
lulesh-comm.cc	623	: : : : : : MPI_Isend(0:26)	13.13us	744.0B	1.66
lulesh-comm.cc	837	: : : : : : MPI_Isend(0:31)	9.48us	24.0B	0.84
lulesh-comm.cc	843	: : : : : : MPI_Waitall(0:)	541.3us	-	0.43
lulesh.cc	1161	: : : : : CommSBN()	-	-	-
lulesh-comm.cc	889	: : : : : : MPI_Comm_rank(0:)	9.32us	-	0.49
lulesh-comm.cc	911	: : : : : : MPI_Wait(0:)	9.25us	-	0.6
lulesh-comm.cc	945	: : : : : : MPI_Wait(0:)	8.6us	-	2.48
lulesh-comm.cc	980	: : : : : : MPI_Wait(0:)	8.52us	-	2.35
lulesh-comm.cc	1039	: : : : : : MPI_Wait(0:)	8.41us	-	1.72
lulesh-comm.cc	1053	: : : : : : MPI_Wait(0:)	8.19us	-	1.64
lulesh-comm.cc	1067	: : : : : : MPI_Wait(0:)	8.31us	-	1.98
lulesh-comm.cc	1251	: : : : : : MPI_Wait(0:)	58.99us	-	1.44
lulesh.cc	1267	: : : : CommRecv()	-	-	-
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	7.77us	-	0.73
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	9.03us	45.05KB	0.79
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	7.37us	45.05KB	1.02
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	7.34us	45.05KB	0.95
lulesh-comm.cc	192	: : : : : MPI_Irecv(0:6)	7.6us	1.45KB	1.12
lulesh-comm.cc	201	: : : : : MPI_Irecv(0:30)	7.51us	1.45KB	1.33
lulesh-comm.cc	210	: : : : : MPI_Irecv(0:26)	7.42us	1.45KB	1.13
lulesh-comm.cc	345	: : : : : MPI_Irecv(0:31)	7.5us	48.0B	0.95
lulesh.cc	1289	: : : : CommSend()	-	-	-
[~ computation ~]		6.15ms	-	2.0	
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	10.26us	-	0.35
lulesh.cc	1292	: : : : CommSyncPosVel()	-	-	-
lulesh-comm.cc	1310	: : : : : MPI_Comm_rank(0:)	22.14us	-	0.73
lulesh-comm.cc	1332	: : : : : MPI_Wait(0:)	274.73us	-	0.79
lulesh-comm.cc	1366	: : : : : MPI_Wait(0:)	72.95us	-	2.46
lulesh-comm.cc	1402	: : : : : MPI_Wait(0:)	23.24us	-	2.3
lulesh-comm.cc	1461	: : : : : MPI_Wait(0:)	198.36us	-	1.79
lulesh-comm.cc	1475	: : : : : MPI_Wait(0:)	7.83us	-	1.83
lulesh-comm.cc	1489	: : : : : MPI_Wait(0:)	7.75us	-	2.08
lulesh-comm.cc	1674	: : : : : MPI_Wait(0:)	333.07us	-	1.73
lulesh.cc	2656	: : : LagrangeElements()	-	-	-
lulesh.cc	2461	: : : : CalcQForElems()	-	-	-
[~ computation ~]		29.32ms	-	2.25	
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	12.61us	21.09KB	0.43
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	9.23us	21.09KB	1.01
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	8.79us	21.09KB	1.17
[~ computation ~]		15.35ms	-	1.91	
lulesh-comm.cc	436	: : : : : MPI_Isend(0:25)	13.14us	21.09KB	1.9
lulesh-comm.cc	477	: : : : : MPI_Isend(0:5)	10.28us	21.09KB	1.89
lulesh-comm.cc	518	: : : : : MPI_Isend(0:1)	10.57us	21.09KB	1.26
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	291.68us	-	0.45
lulesh.cc	2014	: : : : CommMonoQ()	-	-	-
lulesh-comm.cc	1734	: : : : : MPI_Comm_rank(0:)	9.51us	-	0.45
lulesh-comm.cc	1757	: : : : : MPI_Wait(0:)	8.88us	-	0.81
lulesh-comm.cc	1791	: : : : : MPI_Wait(0:)	8.78us	-	2.04
lulesh-comm.cc	1824	: : : : : MPI_Wait(0:)	8.53us	-	2.13
		: END LOOP	-	-	-

Validation

Lulesh 2.0

Execution:

- 6014 code lines (C++) (142 MPI calls)
- 125 parallel processes
- $\approx 160MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 6ms$

Results:

- 30 iterations loop
- 5 long CPU burst
- 1 Data conditioned MPI

```
TimeIncrement()
: [~ computation ~]
: 96.6% => MPI_Allreduce(CommId:1)
LagrangeLeapFrog()
```

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	
lulesh.cc					
lulesh.cc					**2.56**
lulesh.cc					1.57
lulesh.cc					
lulesh.cc					
lulesh.cc					
lulesh.cc					
lulesh-comm.c					0.28
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	10.51us	22.52KB	0.55
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	8.67us	22.52KB	0.96
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	8.43us	22.52KB	1.08
lulesh-comm.cc	192	: : : : : MPI_Irecv(0:6)	8.8us	744.0B	1.08
lulesh-comm.cc	201	: : : : : MPI_Irecv(0:30)	8.54us	744.0B	1.25
lulesh-comm.cc	210	: : : : : MPI_Irecv(0:26)	8.29us	744.0B	1.25
lulesh-comm.cc	345	: : : : : MPI_Irecv(0:31)	8.32us	24.0B	0.9
lulesh.cc	1158	: : : : CommSend()	-	-	
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh-comm.cc	436	: : : : : MPI_Isend(0:25)	14.95us	22.52KB	1.9
lulesh-comm.cc	477	: : : : : MPI_Isend(0:5)	10.07us	22.52KB	1.65
lulesh-comm.cc	518	: : : : : MPI_Isend(0:1)	10.5us	22.52KB	0.93
lulesh-comm.cc	589	: : : : : MPI_Isend(0:6)	14.79us	744.0B	1.68
lulesh-comm.cc	606	: : : : : MPI_Isend(0:30)	12.45us	744.0B	1.86
lulesh-comm.cc	623	: : : : : MPI_Isend(0:26)	13.13us	744.0B	1.66
lulesh-comm.cc	837	: : : : : MPI_Isend(0:31)	9.48us	24.0B	0.84
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	541.3us	-	0.43
lulesh.cc	1161	: : : : CommSBN()	-	-	
lulesh-comm.cc	889	: : : : : MPI_Comm_rank(0:)	9.32us	-	0.49
lulesh-comm.cc	911	: : : : : MPI_Wait(0:)	9.25us	-	0.6
lulesh-comm.cc	945	: : : : : MPI_Wait(0:)	8.6us	-	2.48
lulesh-comm.cc	980	: : : : : MPI_Wait(0:)	8.52us	-	2.35
lulesh-comm.cc	1039	: : : : : MPI_Wait(0:)	8.41us	-	1.72
lulesh-comm.cc	1053	: : : : : MPI_Wait(0:)	8.19us	-	1.64
lulesh-comm.cc	1067	: : : : : MPI_Wait(0:)	8.31us	-	1.98
lulesh-comm.cc	1251	: : : : : MPI_Wait(0:)	58.99us	-	1.44
lulesh.cc	1267	: : : : CommRecv()	-	-	
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	7.77us	-	0.73
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	9.03us	45.05KB	0.79
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	7.37us	45.05KB	1.02
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	7.34us	45.05KB	0.95
lulesh-comm.cc	192	: : : : : MPI_Irecv(0:6)	7.6us	1.45KB	1.12
lulesh-comm.cc	201	: : : : : MPI_Irecv(0:30)	7.51us	1.45KB	1.33
lulesh-comm.cc	210	: : : : : MPI_Irecv(0:26)	7.42us	1.45KB	1.13
lulesh-comm.cc	345	: : : : : MPI_Irecv(0:31)	7.5us	48.0B	0.95
lulesh.cc	1289	: : : : CommSend()	-	-	
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	10.26us	-	0.35
lulesh.cc	1292	: : : : CommSyncPosVel()	-	-	
lulesh-comm.cc	1310	: : : : : MPI_Comm_rank(0:)	22.14us	-	0.73
lulesh-comm.cc	1332	: : : : : MPI_Wait(0:)	274.73us	-	0.79
lulesh-comm.cc	1366	: : : : : MPI_Wait(0:)	72.95us	-	2.46
lulesh-comm.cc	1402	: : : : : MPI_Wait(0:)	23.24us	-	2.3
lulesh-comm.cc	1461	: : : : : MPI_Wait(0:)	198.36us	-	1.79
lulesh-comm.cc	1475	: : : : : MPI_Wait(0:)	7.83us	-	1.83
lulesh-comm.cc	1489	: : : : : MPI_Wait(0:)	7.75us	-	2.08
lulesh-comm.cc	1674	: : : : : MPI_Wait(0:)	333.07us	-	1.73
lulesh.cc	2656	: : : LagrangeElements()	-	-	
lulesh.cc	2461	: : : CalcQForElems()	-	-	
lulesh.cc	1992	: : : : CommRecv()	-	-	
lulesh-comm.cc	**101**	**: : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	12.61us	21.09KB	0.43
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	9.23us	21.09KB	1.01
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	8.79us	21.09KB	1.17
lulesh.cc	2010	: : : : : CommSend()	-	-	
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	15.24us	-	1.52
lulesh-comm.cc	436	: : : : : MPI_Isend(0:25)	13.14us	21.09KB	1.9
lulesh-comm.cc	477	: : : : : MPI_Isend(0:5)	10.28us	21.09KB	1.89
lulesh-comm.cc	518	: : : : : MPI_Isend(0:1)	10.57us	21.09KB	1.26
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	291.68us	-	0.45
lulesh.cc	2014	: : : : CommMonoQ()	-	-	
lulesh-comm.cc	1734	: : : : : MPI_Comm_rank(0:)	9.51us	-	0.45
lulesh-comm.cc	1757	: : : : : MPI_Wait(0:)	8.88us	-	0.81
lulesh-comm.cc	1791	: : : : : MPI_Wait(0:)	8.78us	-	2.04
lulesh-comm.cc	1824	: : : : : MPI_Wait(0:)	8.53us	-	2.13
		: END LOOP	-	-	

Validation

Lulesh 2.0

Execution:

- 6014 code lines (C++) (142 MPI calls)
- 125 parallel processes
- $\approx 160MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 6ms$

Results:

- 30 iterations loop
- 5 long CPU burst
- 1 Data conditioned MPI
- 56 no conditioned MPI

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 30 [id=0.0]	-	-	-
lulesh.cc	2773	: : TimeIncrement()	-	-	-
lulesh.cc	**214**	**: : [~ computation ~]**	**56.63ms**	**-**	**2.56**
lulesh.cc	214	: : 96.6% => MPI_Allreduce(CommId:1)	35.32us	8.0B/8.0B	1.57
lulesh.cc	2774	: : LagrangeLeapFrog()	-	-	-
lulesh.cc	2648	: : LagrangeNodal()	-	-	-
lulesh.cc	1263	: : : CalcForceForNodes()	-	-	-
lulesh.cc	1137	: : : : CommRecv()	-	-	-
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	10.27us	-	0.28
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	10.51us	22.52KB	0.55
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	8.67us	22.52KB	0.96
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	8.43us	22.52KB	1.08
lulesh-comm.cc	192	: : : : : MPI_Irecv(0:6)	8.8us	744.0B	1.08
lulesh-comm.cc	201	: : : : : MPI_Irecv(0:30)	8.54us	744.0B	1.25
lulesh-comm.cc	210	: : : : : MPI_Irecv(0:26)	8.29us	744.0B	1.25
lulesh-comm.cc	345	: : : : : MPI_Irecv(0:31)	8.32us	24.0B	0.9
lulesh.cc	1158	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh-comm.cc	436	: : : : : MPI_Isend(0:25)	14.95us	22.52KB	1.9
lulesh-comm.cc	477	: : : : : MPI_Isend(0:5)	10.07us	22.52KB	1.65
lulesh-comm.cc	518	: : : : : MPI_Isend(0:1)	10.5us	22.52KB	0.93
lulesh-comm.cc	589	: : : : : MPI_Isend(0:6)	14.79us	744.0B	1.68
lulesh-comm.cc	606	: : : : : MPI_Isend(0:30)	12.45us	744.0B	1.86
lulesh-comm.cc	623	: : : : : MPI_Isend(0:26)	13.13us	744.0B	1.66
lulesh-comm.cc	837	: : : : : MPI_Isend(0:31)	9.48us	24.0B	0.84
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	541.3us	-	0.43
lulesh.cc	1161	: : : : CommSBN()	-	-	-
lulesh-comm.cc	889	: : : : : MPI_Comm_rank(0:)	9.32us	-	0.49
lulesh-comm.cc	911	: : : : : MPI_Wait(0:)	9.25us	-	0.6
lulesh-comm.cc	945	: : : : : MPI_Wait(0:)	8.6us	-	2.48
lulesh-comm.cc	980	: : : : : MPI_Wait(0:)	8.52us	-	2.35
lulesh-comm.cc	1039	: : : : : MPI_Wait(0:)	8.41us	-	1.72
lulesh-comm.cc	1053	: : : : : MPI_Wait(0:)	8.19us	-	1.64
lulesh-comm.cc	1067	: : : : : MPI_Wait(0:)	8.31us	-	1.98
lulesh-comm.cc	1251	: : : : : MPI_Wait(0:)	58.99us	-	1.44
lulesh.cc	1267	: : : : CommRecv()	-	-	-
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	7.77us	-	0.73
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	9.03us	45.05KB	0.79
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	7.37us	45.05KB	1.02
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	7.34us	45.05KB	0.95
lulesh-comm.cc	192	: : : : : MPI_Irecv(0:6)	7.6us	1.45KB	1.12
lulesh-comm.cc	201	: : : : : MPI_Irecv(0:30)	7.51us	1.45KB	1.33
lulesh-comm.cc	210	: : : : : MPI_Irecv(0:26)	7.42us	1.45KB	1.13
lulesh-comm.cc	345	: : : : : MPI_Irecv(0:31)	7.5us	48.0B	0.95
lulesh.cc	1289	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	10.26us	-	0.35
lulesh.cc	1292	: : : : CommSyncPosVel()	-	-	-
lulesh-comm.cc	1310	: : : : : MPI_Comm_rank(0:)	22.14us	-	0.73
lulesh-comm.cc	1332	: : : : : MPI_Wait(0:)	274.73us	-	0.79
lulesh-comm.cc	1366	: : : : : MPI_Wait(0:)	72.95us	-	2.46
lulesh-comm.cc	1402	: : : : : MPI_Wait(0:)	23.24us	-	2.3
lulesh-comm.cc	1461	: : : : : MPI_Wait(0:)	198.36us	-	1.79
lulesh-comm.cc	1475	: : : : : MPI_Wait(0:)	7.83us	-	1.83
lulesh-comm.cc	1489	: : : : : MPI_Wait(0:)	7.75us	-	2.08
lulesh-comm.cc	1674	: : : : : MPI_Wait(0:)	333.07us	-	1.73
lulesh.cc	2656	: : : LagrangeElements()	-	-	-
lulesh.cc	2461	: : : : CalcQForElems()	-	-	-
lulesh.cc	1992	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh-comm.cc	119	: : : : : MPI_Irecv(0:25)	12.61us	21.09KB	0.43
lulesh-comm.cc	137	: : : : : MPI_Irecv(0:5)	9.23us	21.09KB	1.01
lulesh-comm.cc	155	: : : : : MPI_Irecv(0:1)	8.79us	21.09KB	1.17
lulesh.cc	2010	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	15.24us	-	1.52
lulesh-comm.cc	436	: : : : : MPI_Isend(0:25)	13.14us	21.09KB	1.9
lulesh-comm.cc	477	: : : : : MPI_Isend(0:5)	10.28us	21.09KB	1.89
lulesh-comm.cc	518	: : : : : MPI_Isend(0:1)	10.57us	21.09KB	1.26
lulesh-comm.cc	843	: : : : : MPI_Waitall(0:)	291.68us	-	0.45
lulesh.cc	2014	: : : : CommMonoQ()	-	-	-
lulesh-comm.cc	1734	: : : : : MPI_Comm_rank(0:)	9.51us	-	0.45
lulesh-comm.cc	1757	: : : : : MPI_Wait(0:)	8.88us	-	0.81
lulesh-comm.cc	1791	: : : : : MPI_Wait(0:)	8.78us	-	2.04
lulesh-comm.cc	1824	: : : : : MPI_Wait(0:)	8.53us	-	2.13
		: END LOOP	-	-	-

Validation

Lulesh 2.0

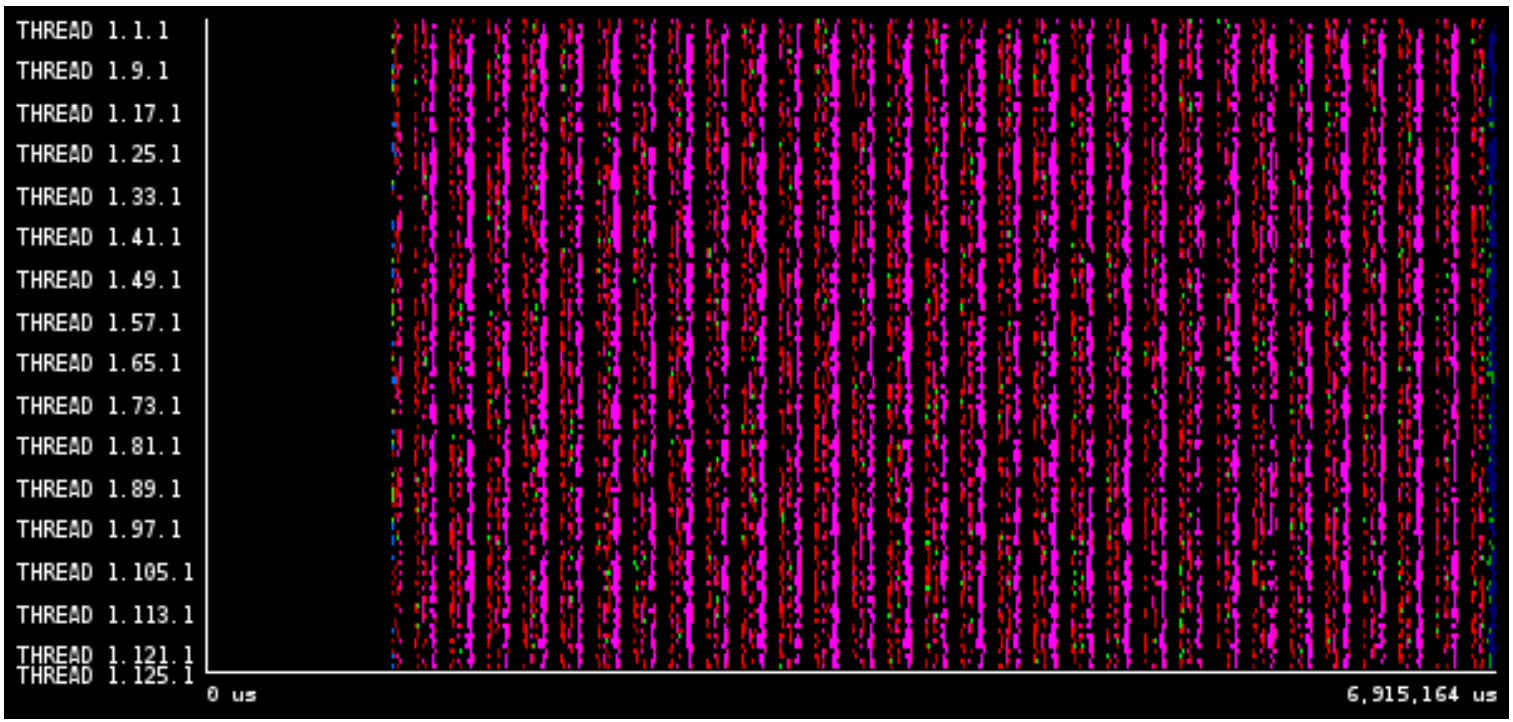

Figure: Lulesh 2.0 125 MPI ranks – 30 iterations

Validation

Lulesh 2.0

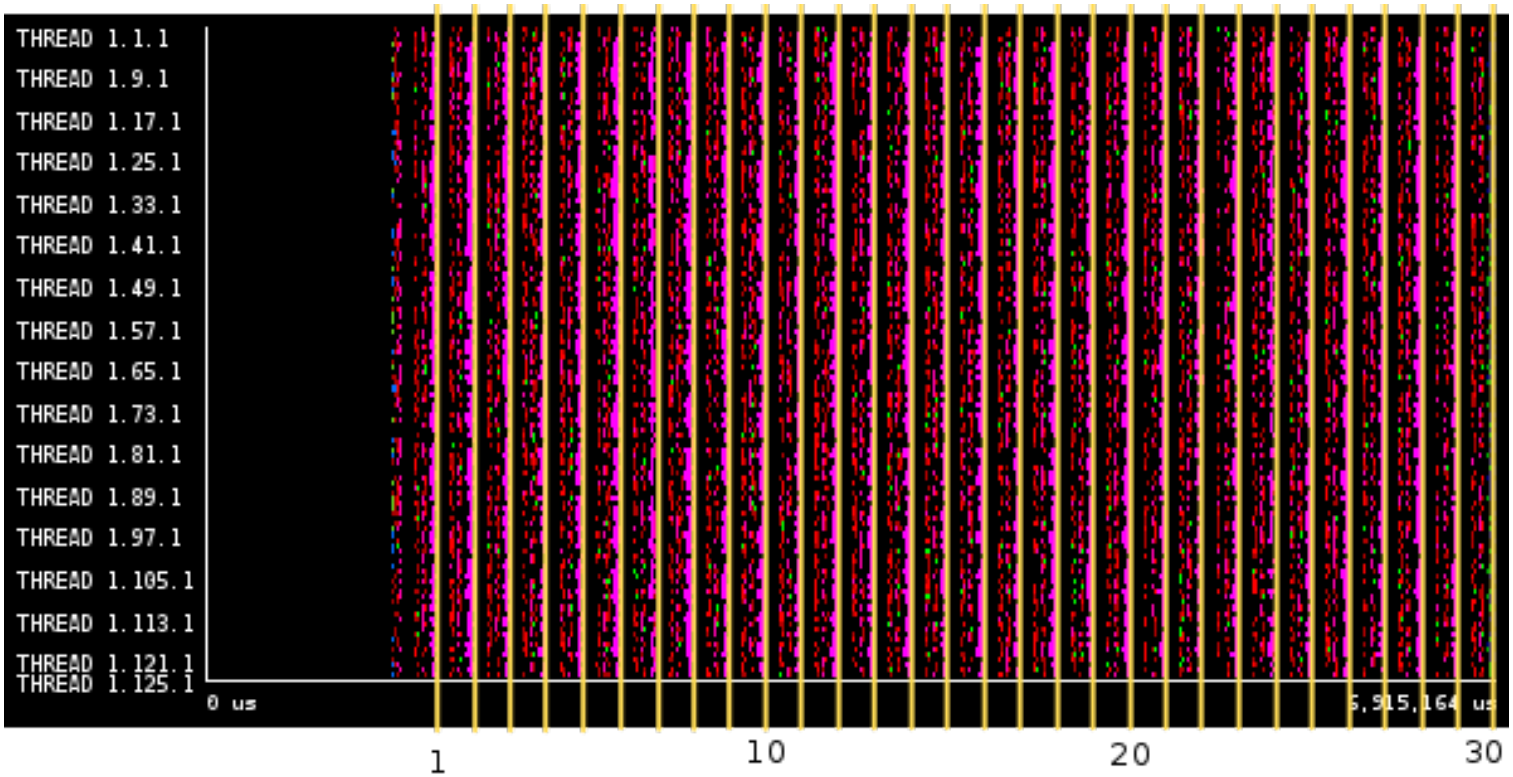

Figure: Lulesh 2.0 125 MPI ranks – 30 iterations

Validation

Lulesh 2.0

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
lulesh.cc	2773	: : TimeIncrement()	-	-	-
lulesh.cc	**214**	**: : : [~ computation ~]**	**56.63ms**	**-**	**2.56**
lulesh.cc	214	: : : 96.6% => MPI_Allreduce(CommId:1)	35.32us	8.0B/8.0B	1.57
lulesh.cc	1158	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh.cc	1289	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh.cc	1992	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh.cc	2010	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	15.24us	-	1.52

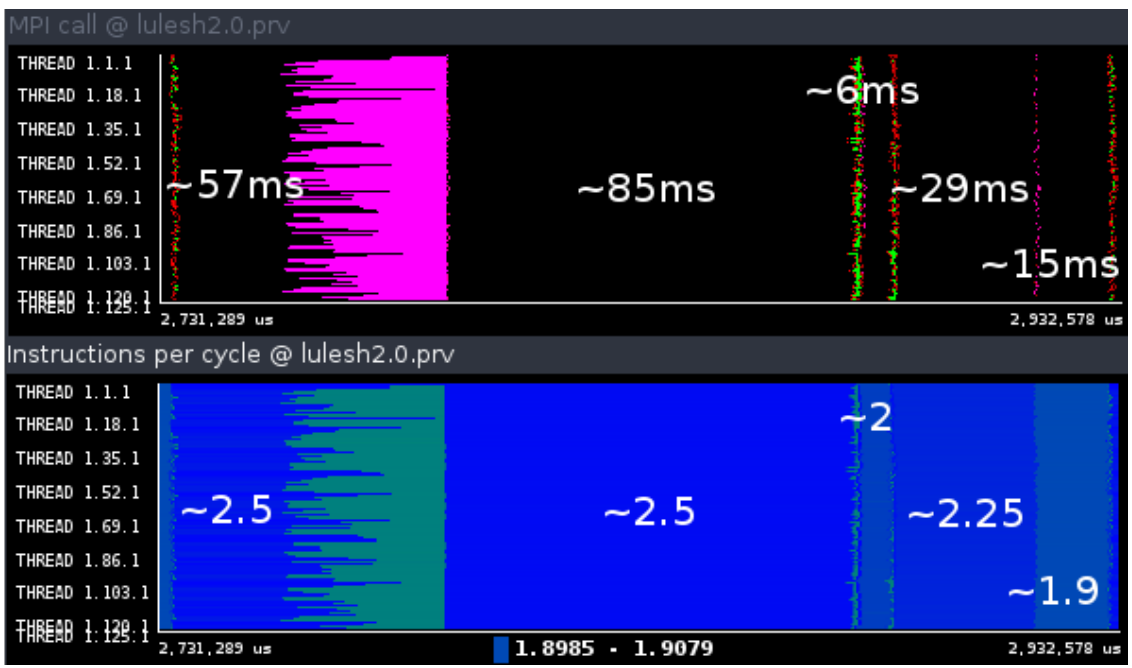

Validation

Lulesh 2.0

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
lulesh.cc	2773	: : TimeIncrement()	-	-	-
lulesh.cc	**214**	**: : : [~ computation ~]**	**56.63ms**	**-**	**2.56**
lulesh.cc	214	: : : 96.6% => MPI_Allreduce(CommId:1)	35.32us	8.0B/8.0B	1.57
lulesh.cc	1158	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh.cc	1289	: : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh.cc	1992	: : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh.cc	2010	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	15.24us	-	1.52

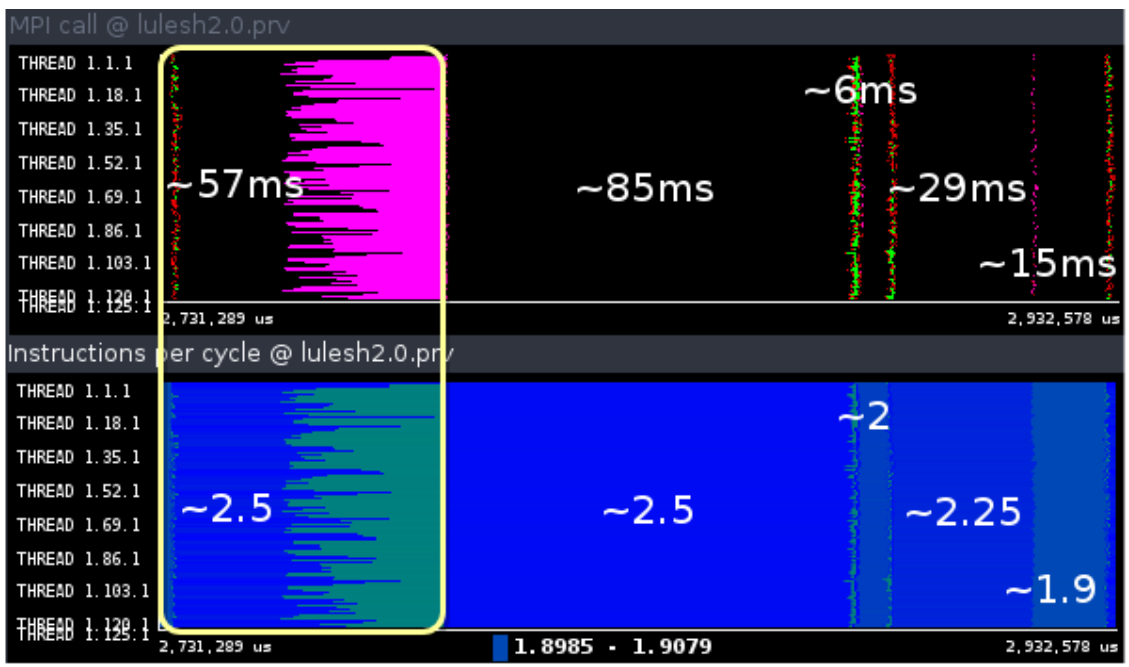

Validation

Lulesh 2.0

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
lulesh.cc	2773	: : TimeIncrement()	-	-	-
lulesh.cc	**214**	**: : : [~ computation ~]**	**56.63ms**	**-**	**2.56**
lulesh.cc	214	: : : 96.6% => MPI_Allreduce(CommId:1)	35.32us	8.0B/8.0B	1.57
lulesh.cc	1158	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh.cc	1289	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh.cc	1992	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh.cc	2010	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	15.24us	-	1.52

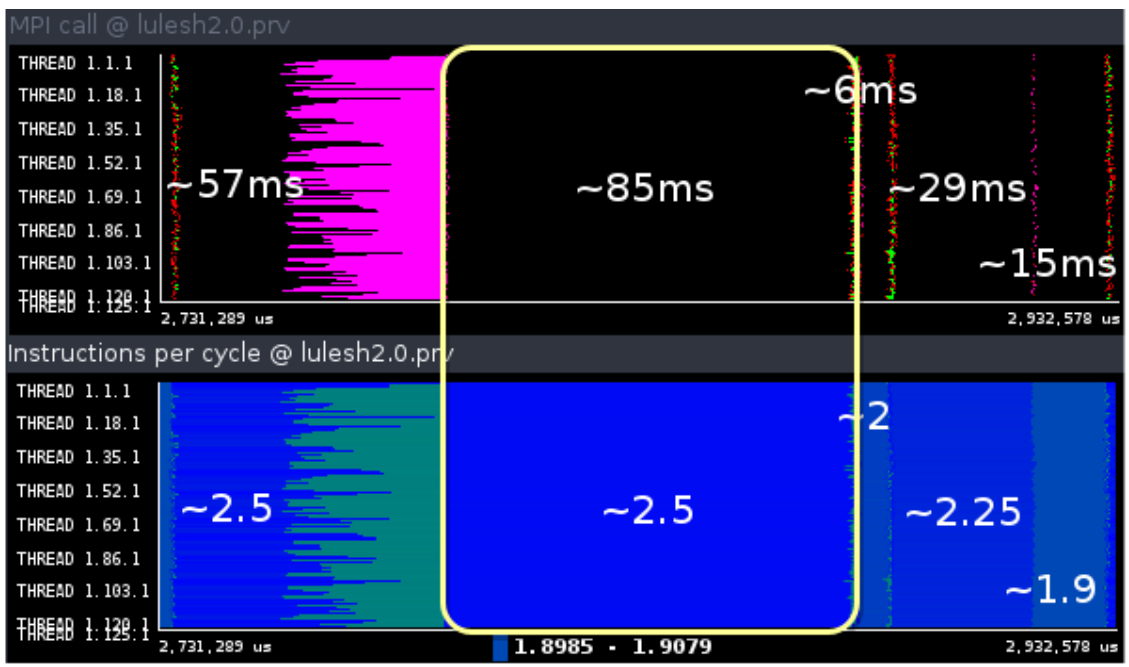

Validation

Lulesh 2.0

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
lulesh.cc	2773	: : TimeIncrement()	-	-	-
lulesh.cc	**214**	**: : : [~ computation ~]**	**56.63ms**	**-**	**2.56**
lulesh.cc	214	: : : 96.6% => MPI_Allreduce(CommId:1)	35.32us	8.0B/8.0B	1.57
lulesh.cc	1158	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh.cc	1289	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh.cc	1992	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh.cc	2010	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	15.24us	-	1.52

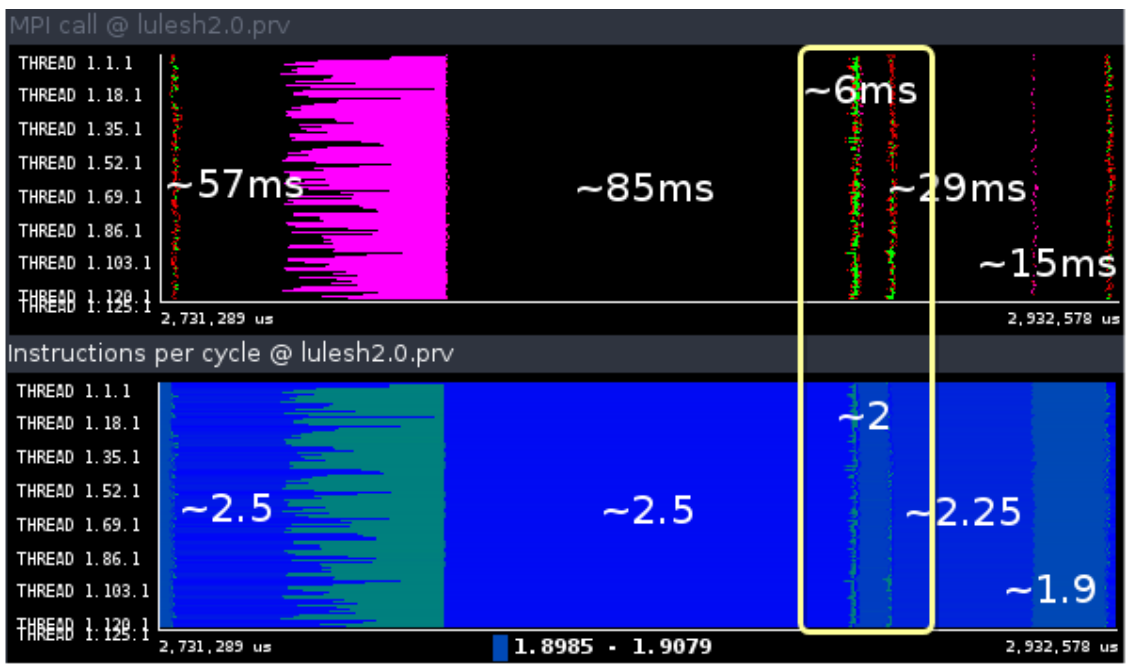

Validation

Lulesh 2.0

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
lulesh.cc	2773	: : TimeIncrement()	-	-	-
lulesh.cc	**214**	**: : : [~ computation ~]**	**56.63ms**	**-**	**2.56**
lulesh.cc	214	: : : 96.6% => MPI_Allreduce(CommId:1)	35.32us	8.0B/8.0B	1.57
lulesh.cc	1158	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh.cc	1289	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh.cc	1992	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh.cc	2010	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	15.24us	-	1.52

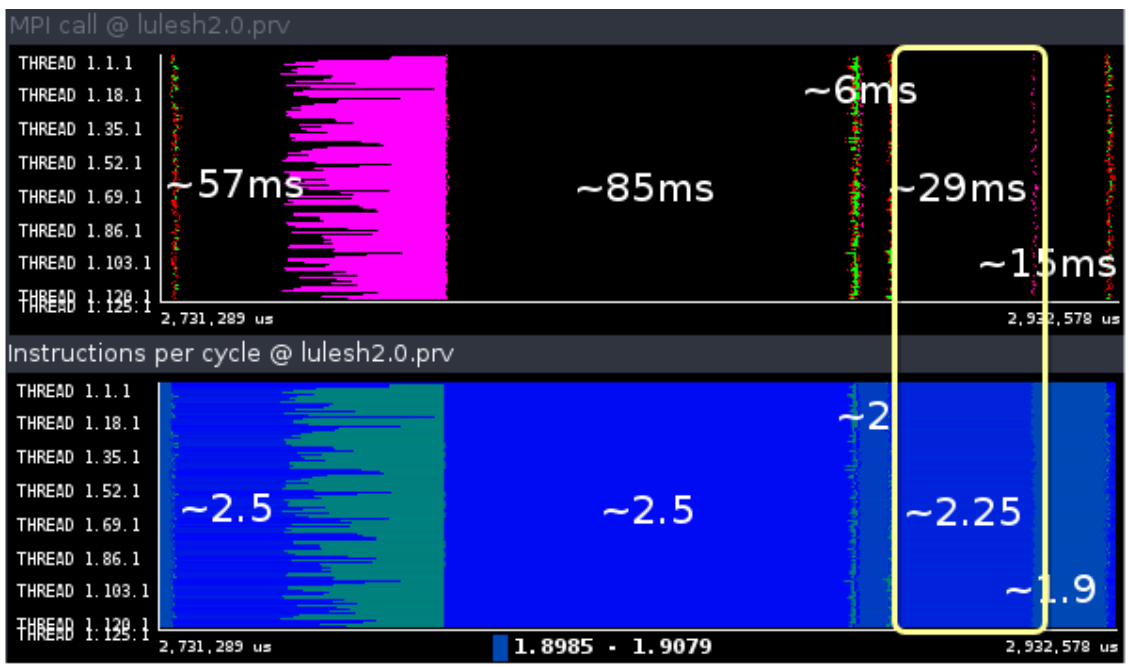

Validation

Lulesh 2.0

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
lulesh.cc	2773	: : TimeIncrement()	-	-	-
lulesh.cc	**214**	**: : : [~ computation ~]**	**56.63ms**	**-**	**2.56**
lulesh.cc	214	: : : 96.6% => MPI_Allreduce(CommId:1)	35.32us	8.0B/8.0B	1.57
lulesh.cc	1158	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**85.2ms**	**-**	**2.48**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	20.08us	-	1.59
lulesh.cc	1289	: : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : [~ computation ~]**	**6.15ms**	**-**	**2.0**
lulesh-comm.cc	401	: : : : : MPI_Comm_rank(0:)	14.58us	-	1.63
lulesh.cc	1992	: : : : : CommRecv()	-	-	-
lulesh-comm.cc	**101**	**: : : : : : [~ computation ~]**	**29.32ms**	**-**	**2.25**
lulesh-comm.cc	101	: : : : : : MPI_Comm_rank(0:)	18.76us	-	1.39
lulesh.cc	2010	: : : : : CommSend()	-	-	-
lulesh-comm.cc	**401**	**: : : : : : [~ computation ~]**	**15.35ms**	**-**	**1.91**
lulesh-comm.cc	401	: : : : : : MPI_Comm_rank(0:)	15.24us	-	1.52

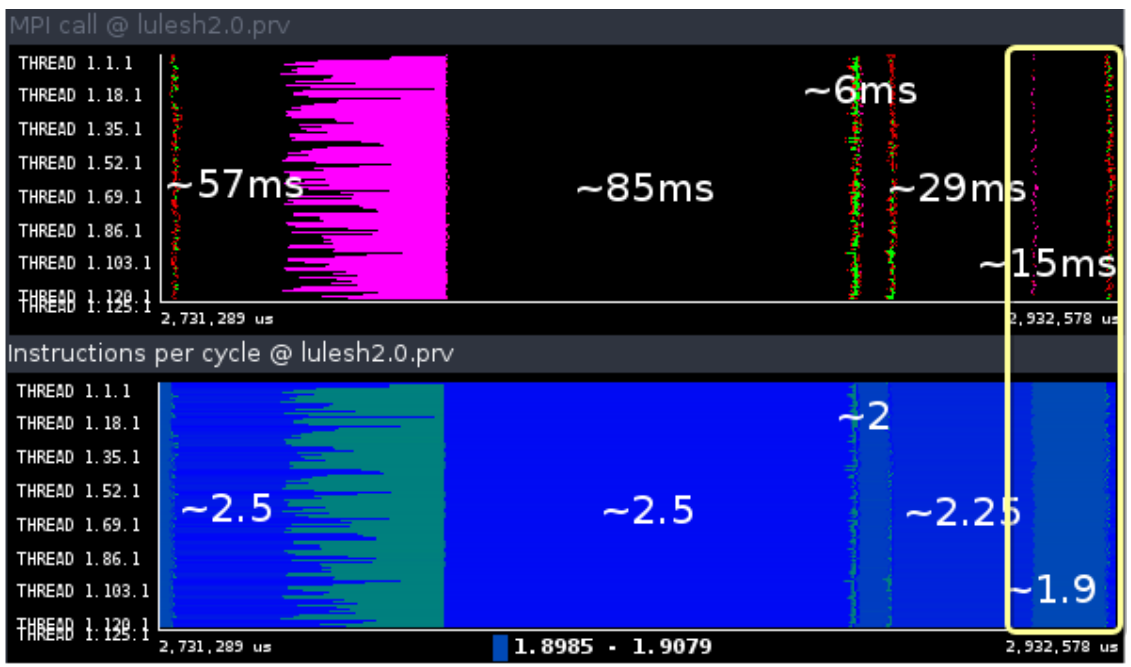

Validation

CG Class A

Execution:

- $\approx$ 2000 code lines (Fortran)
  (40 MPI calls)
- 32 parallel processes
- $\approx 150MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 100us$

Validation

CG Class A

Execution:

- $\approx$ 2000 code lines (Fortran) (40 MPI calls)
- 32 parallel processes
- $\approx 150MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 100us$

Results:

- 15 iterations loop

	0	MAIN__()	-	-	-
		: FOR 1 TO 15	-	-	-
cg.f	475	: : conj_grad_()	-	-	-
		: : : FOR 1 TO 3.0	-	-	-
cg.f	1089	: : : : MPI_Irecv(0:1,2,4)	7.23us	8.0B	1.76
cg.f	1096	: : : : MPI_Send(0:1,2,4)	6.81us	8.0B	1.34
cg.f	1097	: : : : MPI_Wait(0:)	46.79us	-	1.32
		: : : END LOOP	-	-	-
		: : : FOR 1 TO 25.0	-	-	-
		: : : : FOR 1 TO 3.0	-	-	-
cg.f	**1137**	: : : : **[~ computation ~]**	**115.51us**	**-**	**2.32**
cg.f	1137	: : : : MPI_Irecv(0:1,2,4)	7.31us	13.67KB	1.27
cg.f	1144	: : : : MPI_Send(0:1,2,4)	10.52us	13.67KB	1.42
cg.f	1145	: : : : MPI_Wait(0:)	12.07us	-	1.3
		: : : : END LOOP	-	-	-
cg.f	1163	: : : : MPI_Irecv(0:0)	6.7us	13.67KB	2.48
cg.f	1171	: : : : MPI_Send(0:0)	7.04us	13.67KB	1.45
cg.f	1172	: : : : MPI_Wait(0:)	6.45us	-	1.28
		: : : : FOR 1 TO 3.0	-	-	-
cg.f	1207	: : : : : MPI_Irecv(0:1,2,4)	6.55us	8.0B	1.83
cg.f	1214	: : : : : MPI_Send(0:1,2,4)	6.98us	8.0B	1.52
cg.f	1216	: : : : : MPI_Wait(0:)	13.15us	-	1.4
		: : : : END LOOP	-	-	-
		: : : : FOR 1 TO 3.0	-	-	-
cg.f	1262	: : : : : MPI_Irecv(0:1,2,4)	15x	8.0B	2.75
cg.f	1269	: : : : : MPI_Send(0:1,2,4)		8.0B	1.51
cg.f	1270	: : : : : MPI_Wait(0:)		-	1.39
		: : : : END LOOP	-	-	-
		: : : END LOOP	-	-	-
		: : : FOR 1 TO 3.0	-	-	-
cg.f	**1320**	: : : : **[~ computation ~]**	**115.65us**	**-**	**2.31**
cg.f	1320	: : : : MPI_Irecv(0:1,2,4)	6.94us	13.67KB	1.13
cg.f	1327	: : : : MPI_Send(0:1,2,4)	10.55us	13.67KB	1.43
cg.f	1328	: : : : MPI_Wait(0:)	9.14us	-	1.25
		: : : END LOOP	-	-	-
cg.f	1347	: : : MPI_Irecv(0:0)	6.76us	13.67KB	2.47
cg.f	1355	: : : MPI_Send(0:0)	7.13us	13.67KB	1.44
cg.f	1356	: : : MPI_Wait(0:)	6.59us	-	1.19
		: : : FOR 1 TO 3.0	-	-	-
cg.f	1384	: : : : MPI_Irecv(0:1,2,4)	6.57us	8.0B	2.37
cg.f	1391	: : : : MPI_Send(0:1,2,4)	6.63us	8.0B	1.51
cg.f	1392	: : : : MPI_Wait(0:)	25.38us	-	1.38
		: : : END LOOP	-	-	-
		: : FOR 1 TO 3.0	-	-	-
cg.f	499	: : : MPI_Irecv(0:1,2,4)	6.03us	16.0B	3.09
cg.f	506	: : : MPI_Send(0:1,2,4)	6.1us	16.0B	1.53
cg.f	507	: : : MPI_Wait(0:)	6.48us	-	1.35
		: : END LOOP	-	-	-
		: END LOOP	-	-	-

Validation

CG Class A

Execution:

- $\approx$ 2000 code lines (Fortran) (40 MPI calls)
- 32 parallel processes
- $\approx 150MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 100us$

Results:

- 15 iterations loop
- 25 iterations subloop

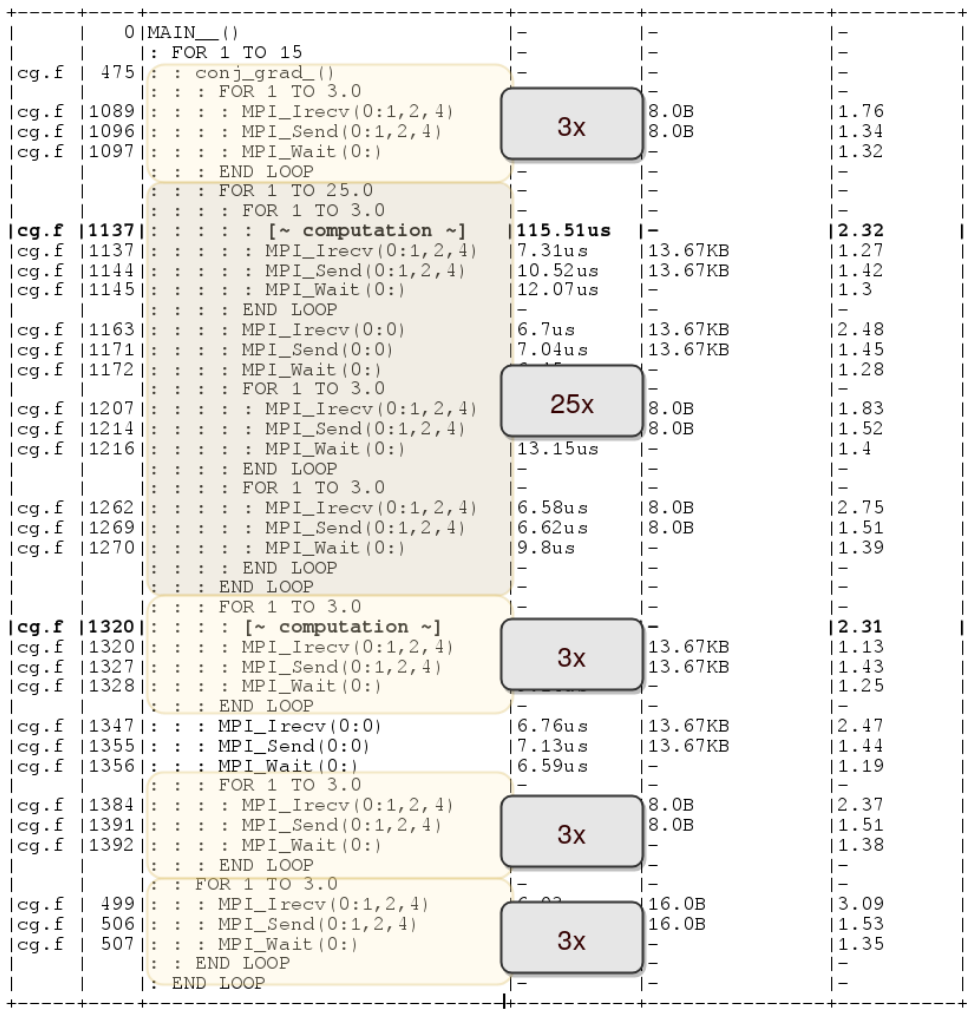

Validation

CG Class A

Execution:

- $\approx$ 2000 code lines (Fortran) (40 MPI calls)
- 32 parallel processes
- $\approx 150MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 100us$

Results:

- 15 iterations loop
- 25 iterations subloop
- Other 3 iterations subloops

```
+-----+----+-----------------------------+----------+--------------+----------+
|     |    0|MAIN__()                     |-         |-             |-         |
|     |     |: FOR 1 TO 15                |-         |-             |-         |
|cg.f | 475|: : conj_grad_()              |-         |-             |-         |
|     |     |: : : FOR 1 TO 3.0           |-         |-             |-         |
|cg.f |1089|: : : : MPI_Irecv(0:1,2,4)    |7.23us    |8.0B          |1.76      |
|cg.f |1096|: : : : MPI_Send(0:1,2,4)     |6.81us    |8.0B          |1.34      |
|cg.f |1097|: : : : MPI_Wait(0:)          |46.79us   |-             |1.32      |
|     |     |: : : END LOOP               |-         |-             |-         |
|     |     |: : : FOR 1 TO 25.0          |-         |-             |-         |
|     |     |: : : : FOR 1 TO 3.0         |-         |-             |-         |
|cg.f |1137|: : : : : [~ computation ~]    |          |-             |2.32      |
|cg.f |1137|: : : : : MPI_Irecv(0:1,2,4)   |   3x     |13.67KB       |1.27      |
|cg.f |1144|: : : : : MPI_Send(0:1,2,4)    |          |13.67KB       |1.42      |
|cg.f |1145|: : : : : MPI_Wait(0:)         |          |-             |1.3       |
|     |     |: : : : END LOOP             |-         |-             |-         |
|cg.f |1163|: : : : MPI_Irecv(0:0)        |6.7us     |13.67KB       |2.48      |
|cg.f |1171|: : : : MPI_Send(0:0)         |7.04us    |13.67KB       |1.45      |
|cg.f |1172|: : : : MPI_Wait(0:)          |6.45us    |-             |1.28      |
|     |     |: : : : FOR 1 TO 3.0         |-         |-             |-         |
|cg.f |1207|: : : : : MPI_Irecv(0:1,2,4)   |   3x     |8.0B          |1.83      |
|cg.f |1214|: : : : : MPI_Send(0:1,2,4)    |          |8.0B          |1.52      |
|cg.f |1216|: : : : : MPI_Wait(0:)         |          |-             |1.4       |
|     |     |: : : : END LOOP             |-         |-             |-         |
|     |     |: : : : FOR 1 TO 3.0         |-         |-             |-         |
|cg.f |1262|: : : : : MPI_Irecv(0:1,2,4)   |   3x     |8.0B          |2.75      |
|cg.f |1269|: : : : : MPI_Send(0:1,2,4)    |          |8.0B          |1.51      |
|cg.f |1270|: : : : : MPI_Wait(0:)         |          |-             |1.39      |
|     |     |: : : : END LOOP             |-         |-             |-         |
|     |     |: : : END LOOP               |-         |-             |-         |
|     |     |: : : FOR 1 TO 3.0           |-         |-             |-         |
|cg.f |1320|: : : : [~ computation ~]      |115.65us  |-             |2.31      |
|cg.f |1320|: : : : MPI_Irecv(0:1,2,4)    |6.94us    |13.67KB       |1.13      |
|cg.f |1327|: : : : MPI_Send(0:1,2,4)     |10.55us   |13.67KB       |1.43      |
|cg.f |1328|: : : : MPI_Wait(0:)          |9.14us    |-             |1.25      |
|     |     |: : : END LOOP               |-         |-             |-         |
|cg.f |1347|: : : MPI_Irecv(0:0)          |6.76us    |13.67KB       |2.47      |
|cg.f |1355|: : : MPI_Send(0:0)           |7.13us    |13.67KB       |1.44      |
|cg.f |1356|: : : MPI_Wait(0:)            |6.59us    |-             |1.19      |
|     |     |: : : FOR 1 TO 3.0           |-         |-             |-         |
|cg.f |1384|: : : : MPI_Irecv(0:1,2,4)    |6.57us    |8.0B          |2.37      |
|cg.f |1391|: : : : MPI_Send(0:1,2,4)     |6.63us    |8.0B          |1.51      |
|cg.f |1392|: : : : MPI_Wait(0:)          |25.38us   |-             |1.38      |
|     |     |: : : END LOOP               |-         |-             |-         |
|     |     |: : FOR 1 TO 3.0             |-         |-             |-         |
|cg.f | 499|: : : MPI_Irecv(0:1,2,4)      |6.03us    |16.0B         |3.09      |
|cg.f | 506|: : : MPI_Send(0:1,2,4)       |6.1us     |16.0B         |1.53      |
|cg.f | 507|: : : MPI_Wait(0:)            |6.48us    |-             |1.35      |
|     |     |: : END LOOP                 |-         |-             |-         |
|     |     |: END LOOP                   |-         |-             |-         |
+-----+----+-----------------------------+----------+--------------+----------+
```

Validation

CG Class A

Execution:

- $\approx$ 2000 code lines (Fortran) (40 MPI calls)
- 32 parallel processes
- $\approx 150MB$ tracefile

Analysis:

- Filter all *loops* $< 10\%$
- CPU burst *time* $> 100us$

Results:

- 15 iterations loop
- 25 iterations subloop
- Other 3 iterations subloops
- 2 long CPU burst above 100us

```
+-----+----+-----------------------------+----------+--------------+----------+
|     |    0|MAIN__()                     |-         |-             |-         |
|     |     |: FOR 1 TO 15                |-         |-             |-         |
|cg.f | 475|: : conj_grad_()              |-         |-             |-         |
|     |     |: : : FOR 1 TO 3.0           |-         |-             |-         |
|cg.f |1089|: : : : MPI_Irecv(0:1,2,4)    |7.23us    |8.0B          |1.76      |
|cg.f |1096|: : : : MPI_Send(0:1,2,4)     |6.81us    |8.0B          |1.34      |
|cg.f |1097|: : : : MPI_Wait(0:)          |46.79us   |-             |1.32      |
|     |     |: : : END LOOP               |-         |-             |-         |
|     |     |: : : FOR 1 TO 25.0          |-         |-             |-         |
|     |     |: : : : FOR 1 TO 3.0         |-         |-             |-         |
|cg.f |1137|: : : : : [~ computation ~]   |115.51us  |-             |2.32      |
|cg.f |1137|: : : : : MPI_Irecv(0:1,2,4)  |7.31us    |13.67KB       |1.27      |
|cg.f |1144|: : : : : MPI_Send(0:1,2,4)   |10.52us   |13.67KB       |1.42      |
|cg.f |1145|: : : : : MPI_Wait(0:)        |12.07us   |-             |1.3       |
|     |     |: : : : END LOOP             |-         |-             |-         |
|cg.f |1163|: : : : MPI_Irecv(0:0)        |6.7us     |13.67KB       |2.48      |
|cg.f |1171|: : : : MPI_Send(0:0)         |7.04us    |13.67KB       |1.45      |
|cg.f |1172|: : : : MPI_Wait(0:)          |6.45us    |-             |1.28      |
|     |     |: : : : FOR 1 TO 3.0         |-         |-             |-         |
|cg.f |1207|: : : : : MPI_Irecv(0:1,2,4)  |6.55us    |8.0B          |1.83      |
|cg.f |1214|: : : : : MPI_Send(0:1,2,4)   |6.98us    |8.0B          |1.52      |
|cg.f |1216|: : : : : MPI_Wait(0:)        |13.15us   |-             |1.4       |
|     |     |: : : : END LOOP             |-         |-             |-         |
|     |     |: : : : FOR 1 TO 3.0         |-         |-             |-         |
|cg.f |1262|: : : : : MPI_Irecv(0:1,2,4)  |6.58us    |8.0B          |2.75      |
|cg.f |1269|: : : : : MPI_Send(0:1,2,4)   |6.62us    |8.0B          |1.51      |
|cg.f |1270|: : : : : MPI_Wait(0:)        |9.8us     |-             |1.39      |
|     |     |: : : : END LOOP             |-         |-             |-         |
|     |     |: : : END LOOP               |-         |-             |-         |
|     |     |: : : FOR 1 TO 3.0           |-         |-             |-         |
|cg.f |1320|: : : : [~ computation ~]      |115.65us  |-             |2.31      |
|cg.f |1320|: : : : MPI_Irecv(0:1,2,4)    |6.94us    |13.67KB       |1.13      |
|cg.f |1327|: : : : MPI_Send(0:1,2,4)     |10.55us   |13.67KB       |1.43      |
|cg.f |1328|: : : : MPI_Wait(0:)          |9.14us    |-             |1.25      |
|     |     |: : : END LOOP               |-         |-             |-         |
|cg.f |1347|: : : MPI_Irecv(0:0)          |6.76us    |13.67KB       |2.47      |
|cg.f |1355|: : : MPI_Send(0:0)           |7.13us    |13.67KB       |1.44      |
|cg.f |1356|: : : MPI_Wait(0:)            |6.59us    |-             |1.19      |
|     |     |: : : FOR 1 TO 3.0           |-         |-             |-         |
|cg.f |1384|: : : : MPI_Irecv(0:1,2,4)    |6.57us    |8.0B          |2.37      |
|cg.f |1391|: : : : MPI_Send(0:1,2,4)     |6.63us    |8.0B          |1.51      |
|cg.f |1392|: : : : MPI_Wait(0:)          |25.38us   |-             |1.38      |
|     |     |: : : END LOOP               |-         |-             |-         |
|     |     |: : FOR 1 TO 3.0             |-         |-             |-         |
|cg.f | 499|: : : MPI_Irecv(0:1,2,4)      |6.03us    |16.0B         |3.09      |
|cg.f | 506|: : : MPI_Send(0:1,2,4)       |6.1us     |16.0B         |1.53      |
|cg.f | 507|: : : MPI_Wait(0:)            |6.48us    |-             |1.35      |
|     |     |: : END LOOP                 |-         |-             |-         |
|     |     |: END LOOP                   |-         |-             |-         |
+-----+----+-----------------------------+----------+--------------+----------+
```

Validation

CG Class A

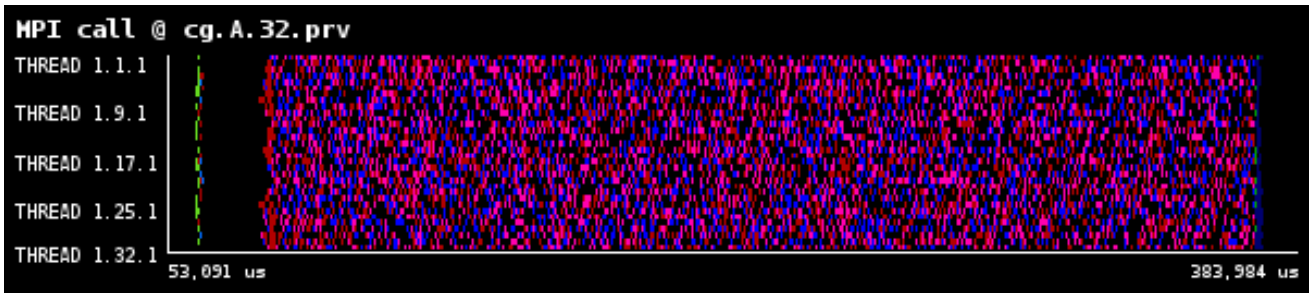

Validation

CG Class A

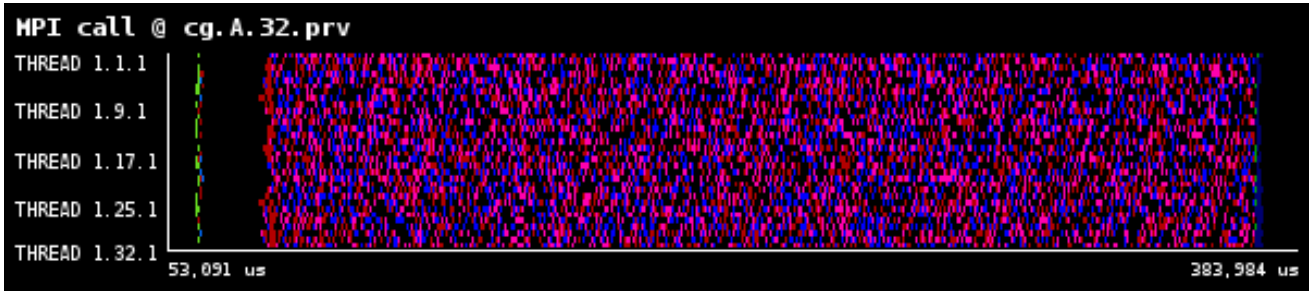

Showing just **MPI_Irecv** on **cg.f : 1089**

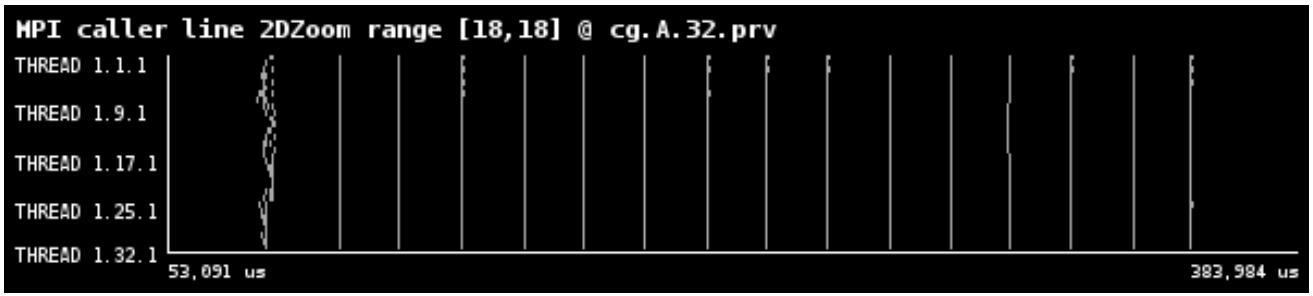

Validation

CG Class A

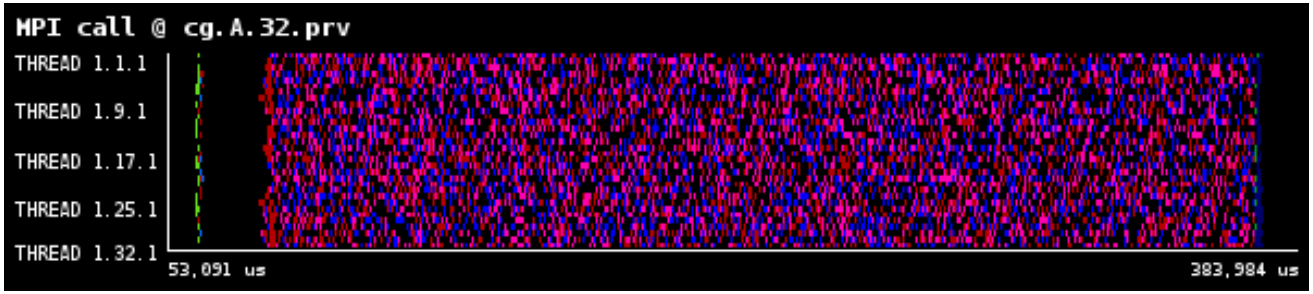

Showing just **MPI_Irecv** on **cg.f : 1089**

Validation

CG Class A

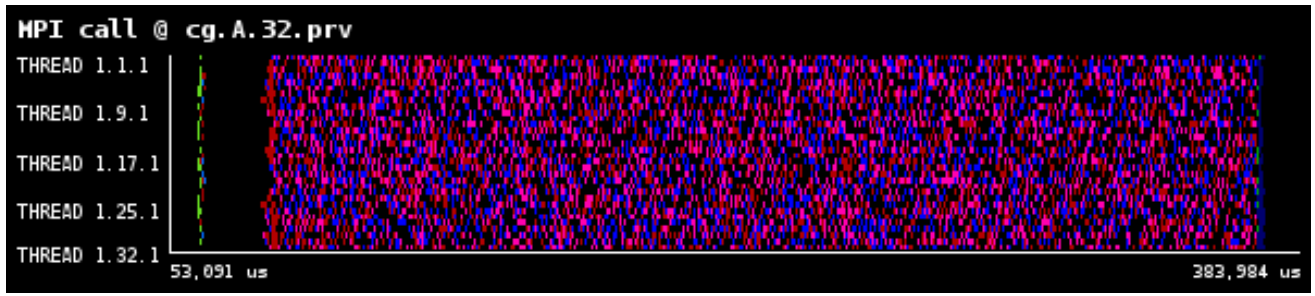

Showing just **MPI_Irecv** on **cg.f : 1089**

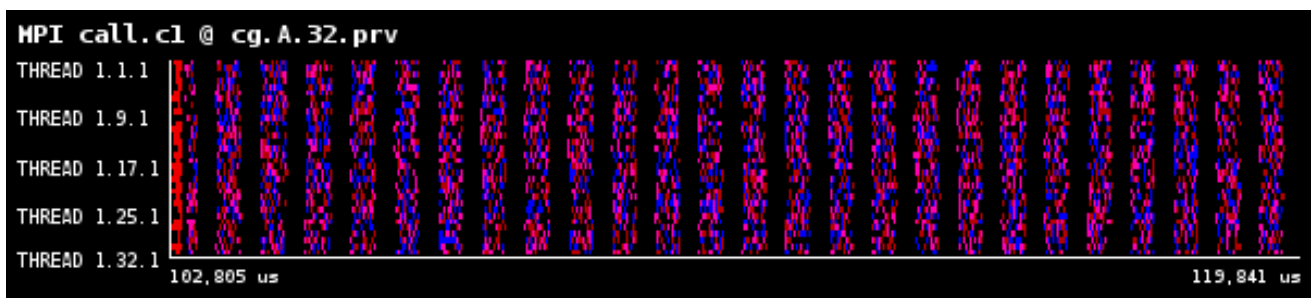

Validation

CG Class A

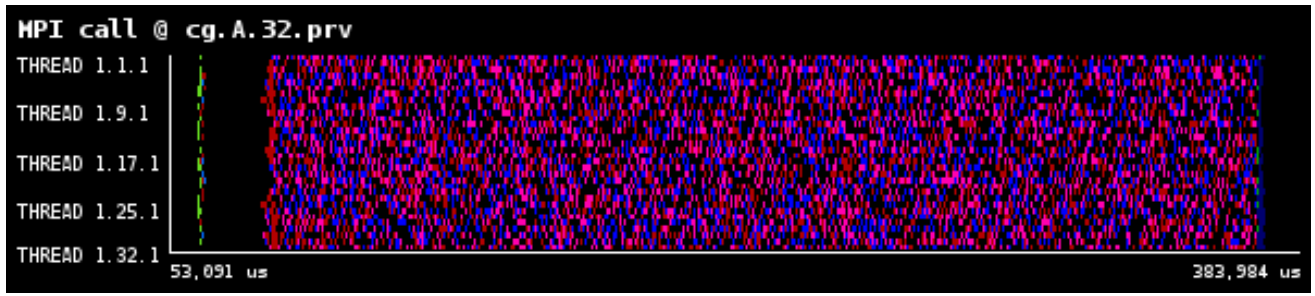

Showing just **MPI_Irecv** on **cg.f : 1089**

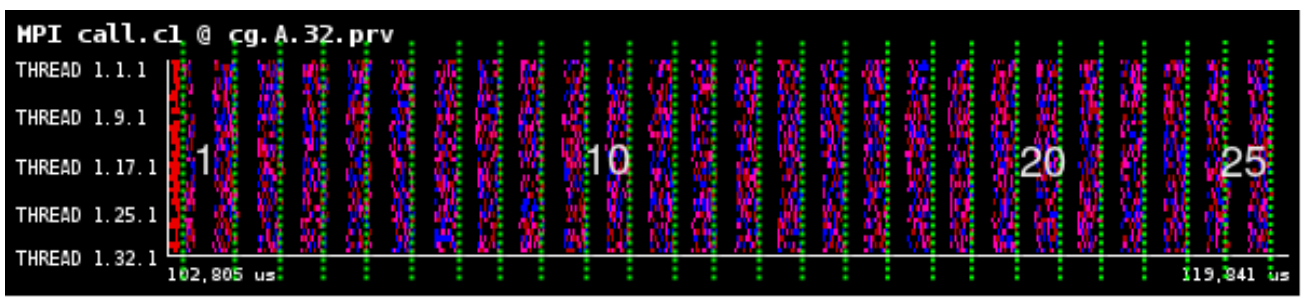

Outline

Scalability

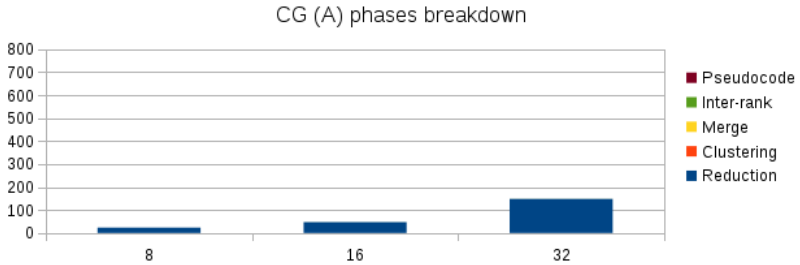

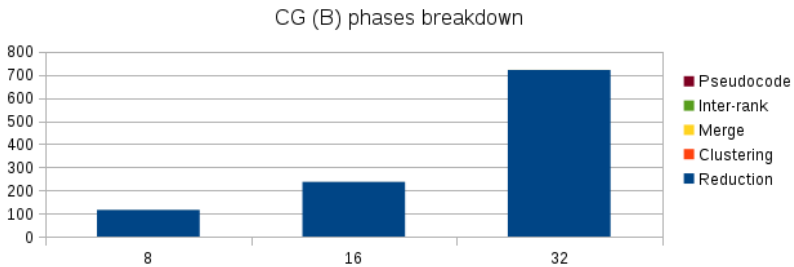

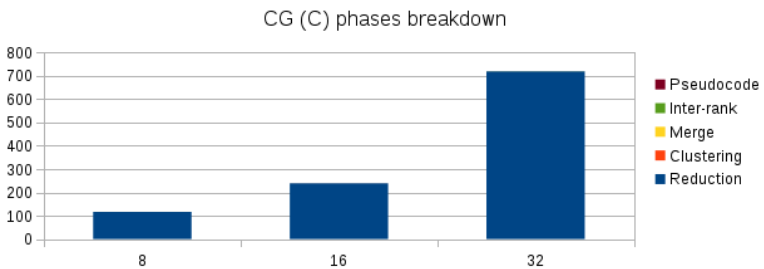

Scalability

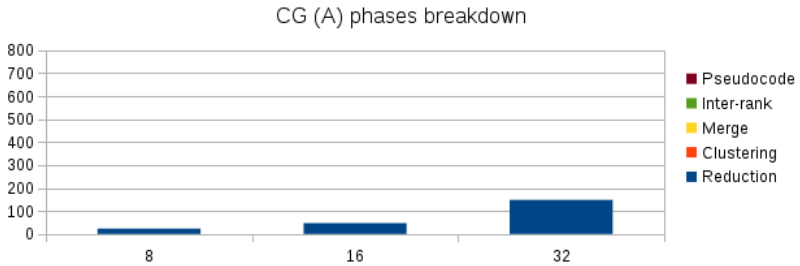

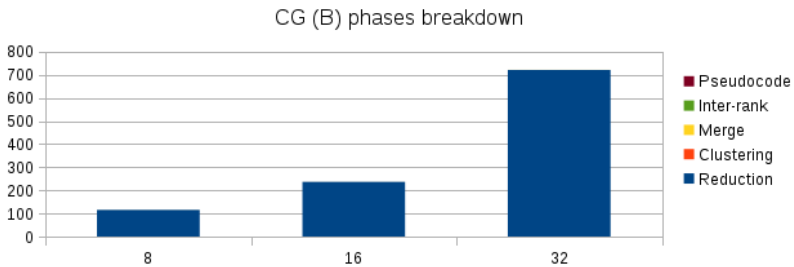

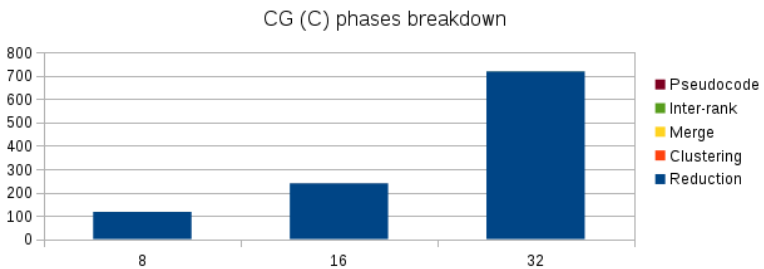

Reduction step is clearly dominating the execution time.

Scalability

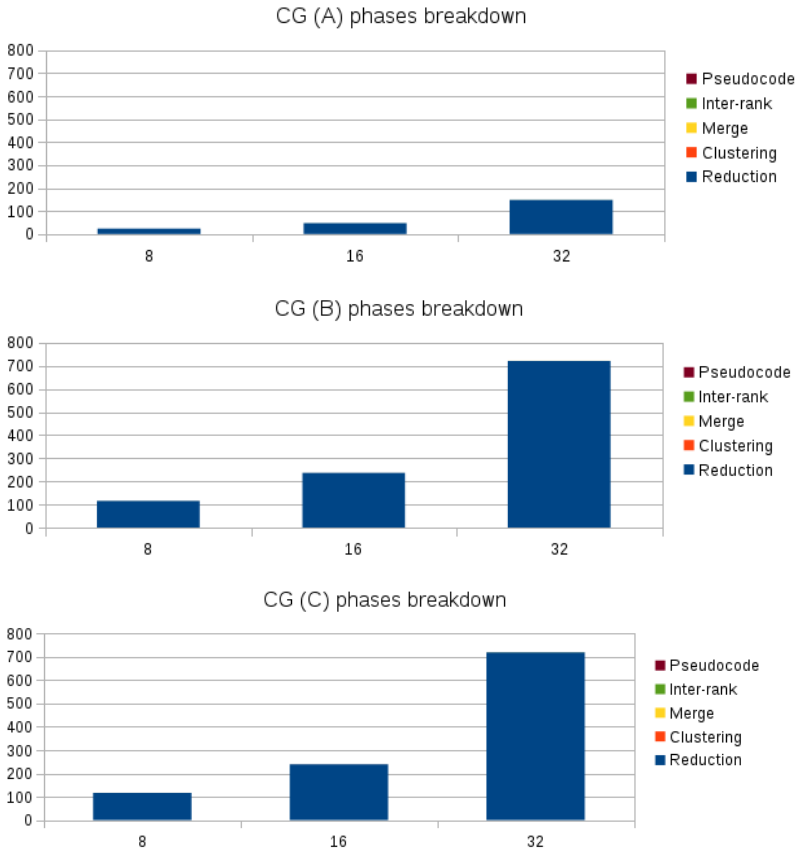

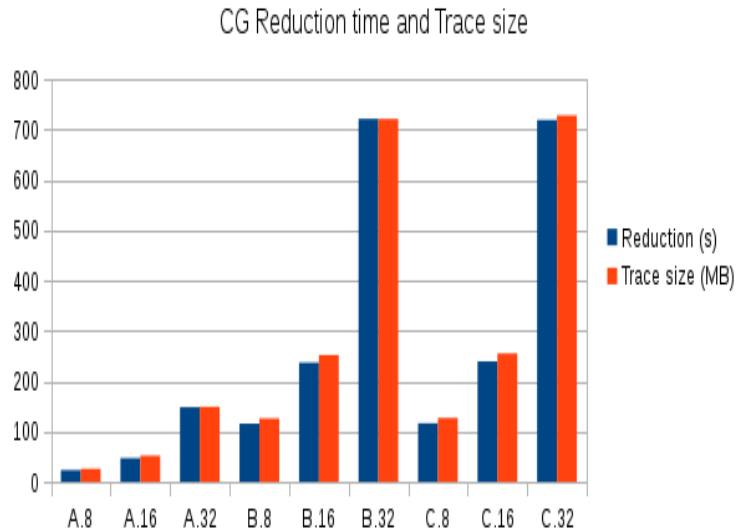

Reduction step is clearly dominating the execution time.

Reduction step time is highly correlated with trace size.

Outline

Conclusions

- New approach for application structure detection has been explored.

Conclusions

- New approach for application structure detection has been explored.
- Have demonstrated that it works well with HPC applications.

Conclusions

- New approach for application structure detection has been explored.
- Have demonstrated that it works well with HPC applications.
- Number of iterations and iterations mean time has been identified as valuable features for loops classification.

Conclusions

- New approach for application structure detection has been explored.
- Have demonstrated that it works well with HPC applications.
- Number of iterations and iterations mean time has been identified as valuable features for loops classification.
- Execution time is clearly dominated by Reduction step that presents linear complexity with the trace size.

Outline

Future work

- Improve memory usage by timestamps compression.

Future work

- Improve memory usage by timestamps compression.
- Improve Reduce phase times by parallelizing the reduction using for example well-known infrastructures like Hadoop.

Future work

- Improve memory usage by timestamps compression.
- Improve Reduce phase times by parallelizing the reduction using for example well-known infrastructures like Hadoop.
- Explore the possibility to use sampling techniques to get more detailed structure of an application.

Master thesis

Master in Research and Innovation

Inferring programs structure from an execution trace

Author

Juan Francisco Martínez Vera

Supervisor

Jesús Labarta Mancho

Facultat d'Informàtica de Barcelona (FIB)

Universitat Politècnica de Catalunya (UPC)

Refinement

Even if it **works pretty well** in general.
Some problems **have been identified**.

1. **Cluster aliasing** when two different loops behaves similarly enough over our defined space
2. **Hidden superloop** when not detected superloop prevents from a good structure detection.
3. **Cluster split** when MPI calls belonging to the same loop behaves in a different way.

All three **have been solved by looking to timestamps**.

Proposal modifications

Cluster aliasing

Keynote

Aliasing can be detected and solved **looking at timestamps**.

for $1 to 100$

do $\begin{cases} MPI\_A() \\ MPI\_B() \end{cases}$

for $1 to 100$

do $\begin{cases} MPI\_C() \\ MPI\_D() \end{cases}$

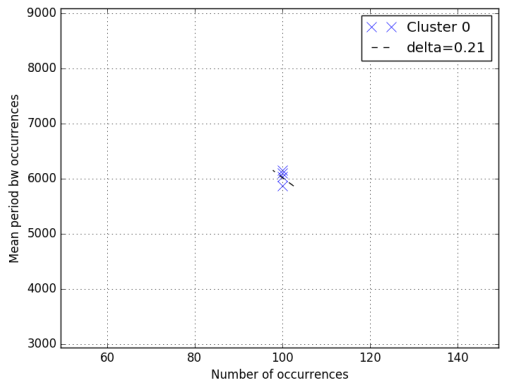

Proposal modifications

Cluster aliasing

Keynote

Aliasing can be detected and solved **looking at timestamps**.

for $1 to 100$
  do $\begin{cases} MPI\_A() \\ MPI\_B() \end{cases}$
for $1 to 100$
  do $\begin{cases} MPI\_C() \\ MPI\_D() \end{cases}$

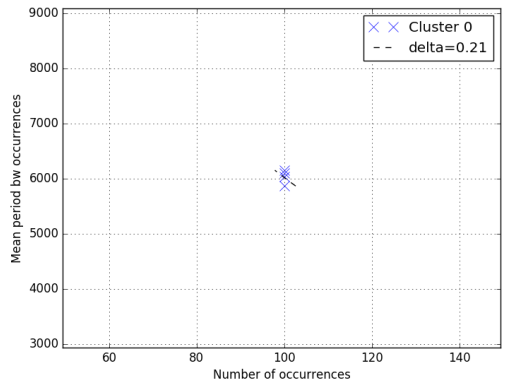

MPI_A	1	3
MPI_B	2	4
MPI_C	5	7
MPI_D	6	8

Proposal modifications

Cluster aliasing

Keynote

Aliasing can be detected and solved **looking at timestamps**.

for $1 to 100$
  do $\begin{cases} MPI\_A() \\ MPI\_B() \end{cases}$
for $1 to 100$
  do $\begin{cases} MPI\_C() \\ MPI\_D() \end{cases}$

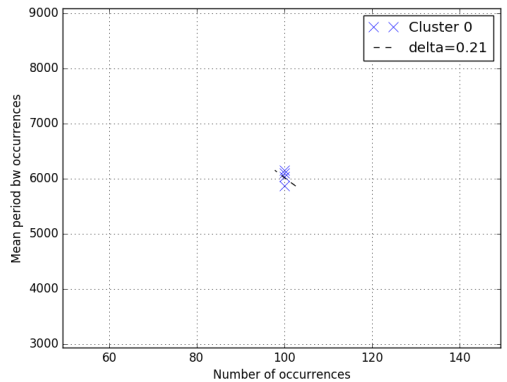

MPI_A	1	3
MPI_B	2	4
MPI_C	5	7
MPI_D	6	8

MPI_A	1	3		
MPI_B	2	4		
MPI_C			5	7
MPI_D			6	8

Proposal modifications

Hidden superloop

Keynote

Similarly than with aliasing, **looking at timestamps**.

for $1 to 100$

do {
- **for** $1 to 2$ **do** { $MPI\_A()$
- **for** $1 to 2$ **do** { $MPI\_B()$

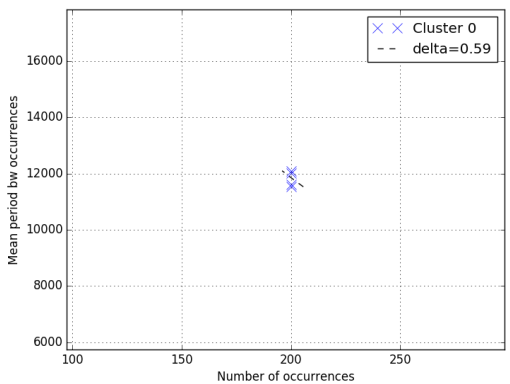

Proposal modifications

Hidden superloop

Keynote

Similarly than with aliasing, **looking at timestamps**.

for $1 to 100$

do { **for** $1 to 2$ **do** $\{MPI\_A()$; **for** $1 to 2$ **do** $\{MPI\_B()$

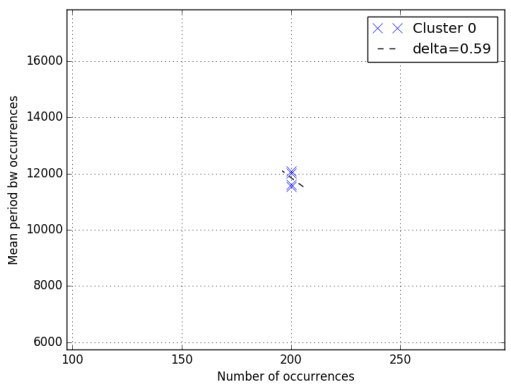

MPI_A	1	2	5	6
MPI_B	3	4	7	8

Proposal modifications

Hidden superloop

Keynote

Similarly than with aliasing, **looking at timestamps**.

for $1 to 100$

do {
- **for** $1 to 2$ **do** $\{MPI\_A()$
- **for** $1 to 2$ **do** $\{MPI\_B()$

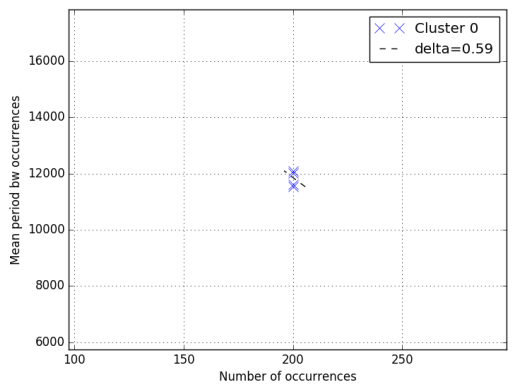

MPI_A	1	2	5	6
MPI_B	3	4	7	8

MPI_A	1	2				5	6	
MPI_B				3	4			7	8

Proposal modifications

Cluster split

$$\textbf{for } i = 1 \text{ to } 10$$
$$\textbf{do} \begin{cases} MPI\_B() \\ \textbf{if } \quad odd(i) \\ \quad \textbf{then } MPI\_A() \\ MPI\_B() \end{cases}$$

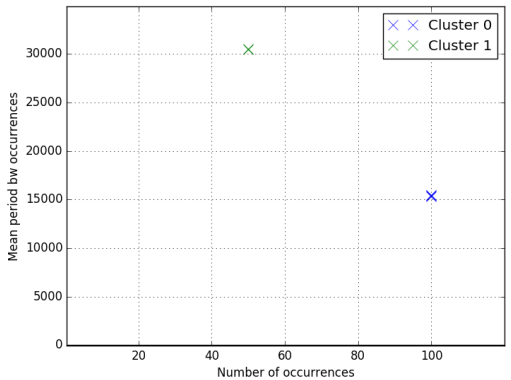

Proposal modifications

Cluster split

$$\textbf{for } i = 1 \text{ to } 10 \textbf{ do } \begin{cases} MPI\_B() \\ \textbf{if } \ odd(i) \\ \quad \textbf{then } MPI\_A() \\ MPI\_B() \end{cases}$$

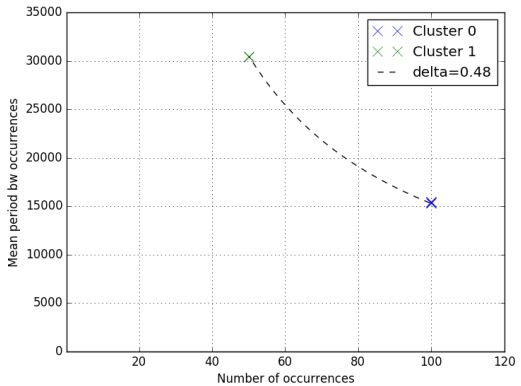

Whether MPI call is conditioned

- #repetitions and mean time bw. reps. change (in same proportion).

Proposal modifications

Cluster split

$$\textbf{for } i = 1 \text{ to } 10 \textbf{ do } \begin{cases} MPI\_B() \\ \textbf{if } odd(i) \\ \quad \textbf{then } MPI\_A() \\ MPI\_B() \end{cases}$$

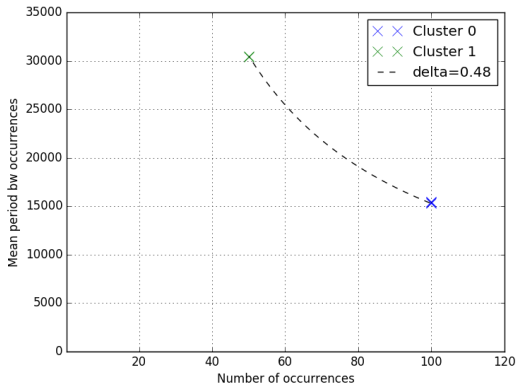

Whether MPI call is conditioned

- #repetitions and mean time bw. reps. change (in same proportion).

By "reverse loop merging" and times checking

FILE	LINE	PSEUDOCODE	E(TIME)	E(SIZE)	E(IPC)
	0	main()	-	-	-
		: FOR 1 TO 100	-	-	-
		: : IF rank in [1]	-	-	-
test-9.c	20	: : : MPI_Send(1:0)	5.01us	4.0B	1.1
		: : IF rank in [0]	-	-	-
test-9.c	22	: : : MPI_Recv()	5.3us	-	1.03
test-9.c	24	: : 50.0% => MPI_Barrier(CommId:1)	5.99us	-	0.97
		: : IF rank in [1]	-	-	-
test-9.c	26	: : : MPI_Send(1:0)	4.9us	4.0B	1.14
		: : IF rank in [0]	-	-	-
test-9.c	28	: : : MPI_Recv()	5.23us	-	1.1
		: END LOOP	-	-	-

Data analysis

Data acquisition

Desired data...

- PCA[2]: Set of observations with set of features, i.e. Set of loops with with **iteration-level aggregated features**.
- Variable Importance Same data as in production traces but labeled, i.e. textbfWith information to what loop every MPI call belongs to.

How to obtain it

1. The needed information is not in trace so ...
2. **Manually inject monitors** to the source code
   - Very tough and error prone for huge source codes
3. Alternatively use source-to-source compiler to **do it utomaticaly**
   - Mercurium

[2] Principal component analysis

Data analysis

Data acquisition

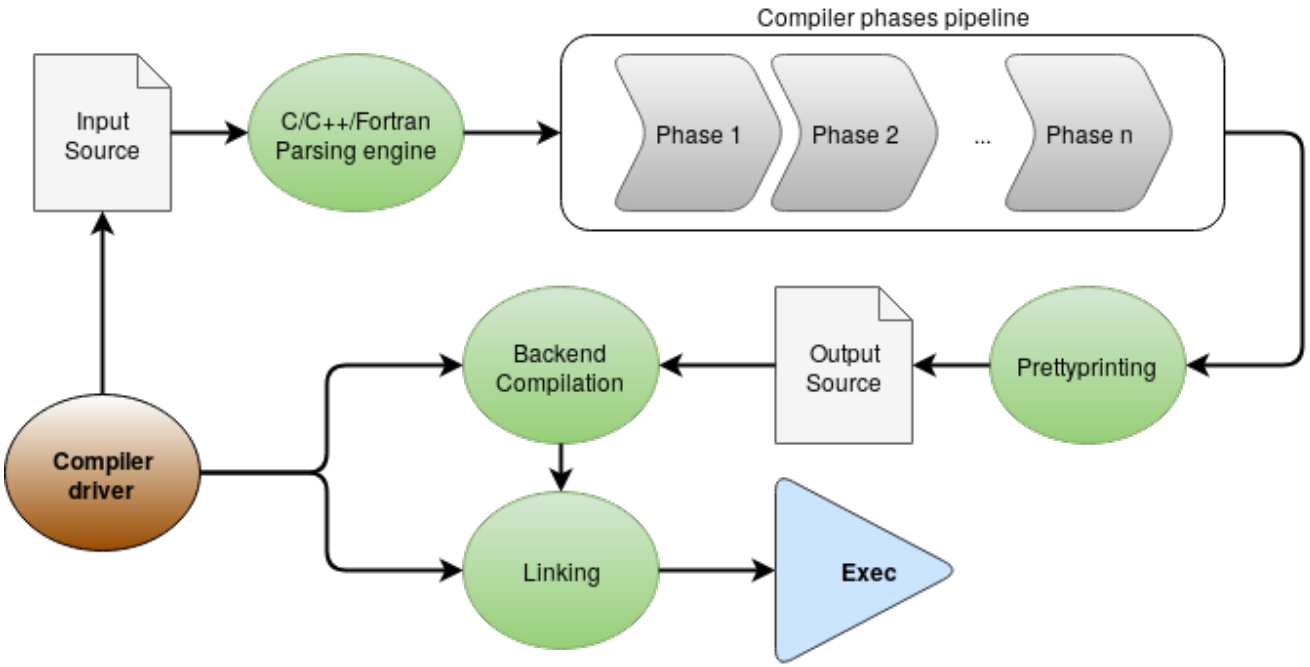

Figure: Mercurium internals overview

Data analysis

Data acquisition

Algorithm 7.1: $\mathrm{PCA}()$

$MLoopInit(loop_{line}, loop_{file})$
for $i \in I$
do $\begin{cases} MIterInit(chance) \\ \ldots \\ MIterFini() \end{cases}$
$MLoopFini()$

- MLoopInit and MLoopFini fire loops boundaries
- MIterInit and MiterFini fire iteration boundaries with hwc information o trace

Data analysis

Data acquisition

Algorithm 7.3: PCA()

MLoopInit($loop_{line}$, $loop_{file}$)
for $i \in I$
do {
  MIterInit(*chance*)
  ...
  MIterFini()
}
MLoopFini()

- MLoopInit and MLoopFini fire loops boundaries
- MIterInit and MiterFini fire iteration boundaries with hwc information o trace

Algorithm 7.4: VarImp()

MLoopInit($loop_{line}$, $loop_{file}$)
for $i \in I$
do {
  ...
  MBefCall()
  MPI_Call()
  MPIAftCall()
  ...
}
MLoopFini()

- MLoopInit and MLoopFini keep track entry/exit loops
- MPIBefCall and MPIAftCall fire loop id

Data analysis

Analysis of results: **PCA**

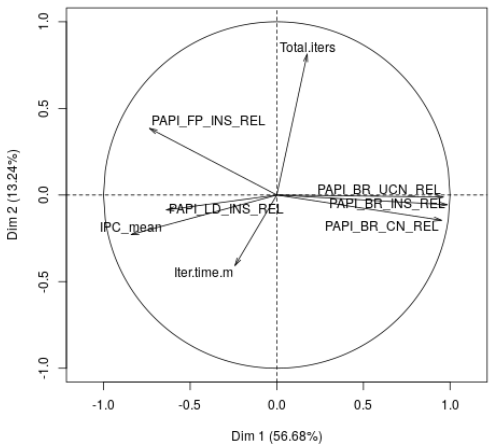

BT

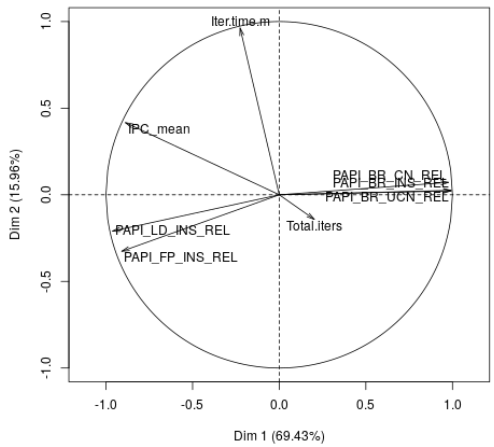

MG

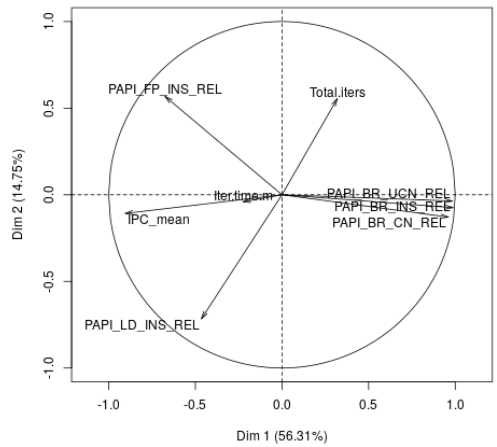

SP

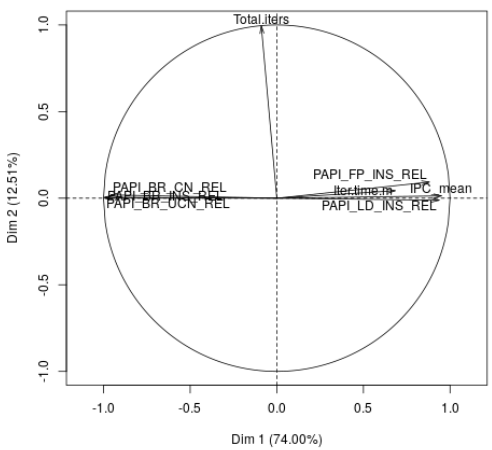

CG

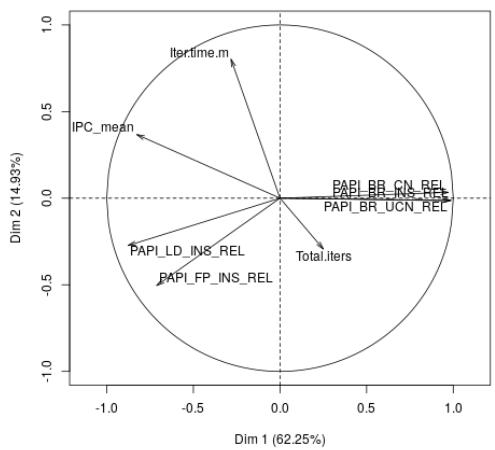

FT

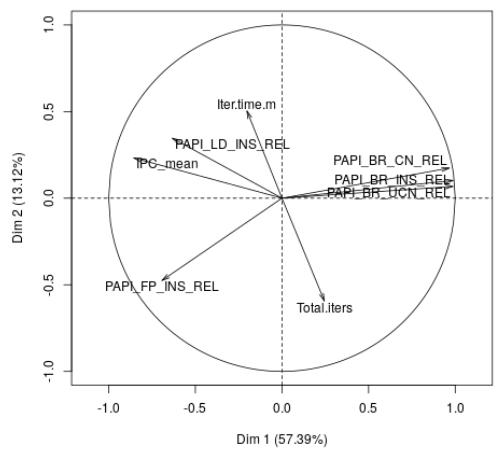

LU

Validation

Lulesh 2.0

```
|lulesh.cc          |2775|: : TimeIncrement()                          |-          |-              |-          |
|lulesh.cc          | 214|: : : [~ computation ~]                      |56.63ms    |-              |2.56       |
|lulesh.cc          | 214|: : : 96.6% => MPI_Allreduce(CommId:1)       |35.32us    |8.0B/8.0B      |1.57       |
|lulesh.cc          |2774|: : LagrangeLeapFrog()                       |-          |-              |-          |
|lulesh.cc          |2648|: : : LagrangeNodal()                        |-          |-              |-          |
|lulesh.cc          |1263|: : : : CalcForceForNodes()                  |-          |-              |-          |
|lulesh.cc          |1137|: : : : : CommRecv()                         |-          |-              |-          |
|lulesh-comm.cc     | 101|: : : : : : MPI_Comm_rank(0:)                |10.27us    |-              |0.28       |
|lulesh-comm.cc     | 119|: : : : : : MPI_Irecv(0:25)                  |10.51us    |22.52KB        |0.55       |
|lulesh-comm.cc     | 137|: : : : : : MPI_Irecv(0:5)                   |8.67us     |22.52KB        |0.96       |
|lulesh-comm.cc     | 155|: : : : : : MPI_Irecv(0:1)                   |8.43us     |22.52KB        |1.08       |
|lulesh-comm.cc     | 192|: : : : : : MPI_Irecv(0:6)                   |8.8us      |744.0B         |1.08       |
|lulesh-comm.cc     | 201|: : : : : : MPI_Irecv(0:30)                  |8.54us     |744.0B         |1.25       |
|lulesh-comm.cc     | 210|: : : : : : MPI_Irecv(0:26)                  |8.29us     |744.0B         |1.25       |
|lulesh-comm.cc     | 345|: : : : : : MPI_Irecv(0:31)                  |8.32us     |24.0B          |0.9        |
|lulesh.cc          |1158|: : : : : CommSend()                         |-          |-              |-          |
|lulesh-comm.cc     | 401|: : : : : : [~ computation ~]                |85.2ms     |-              |2.48       |
|lulesh-comm.cc     | 401|: : : : : : MPI_Comm_rank(0:)                |20.08us    |-              |1.59       |
```

Table: Total MPI calls

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Rank 0	300	510	510	90	29	270

Table: MPI calls counts by communication phases

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Comm phase 1	0	7	0	0	1	1
Comm phase 2	7	7	7	1	0	3
Comm phase 3	0	0	7	1	0	2
Comm phase 4	0	3	0	0	0	1
Comm phase 5	3	0	3	1	0	2

Validation

Lulesh 2.0

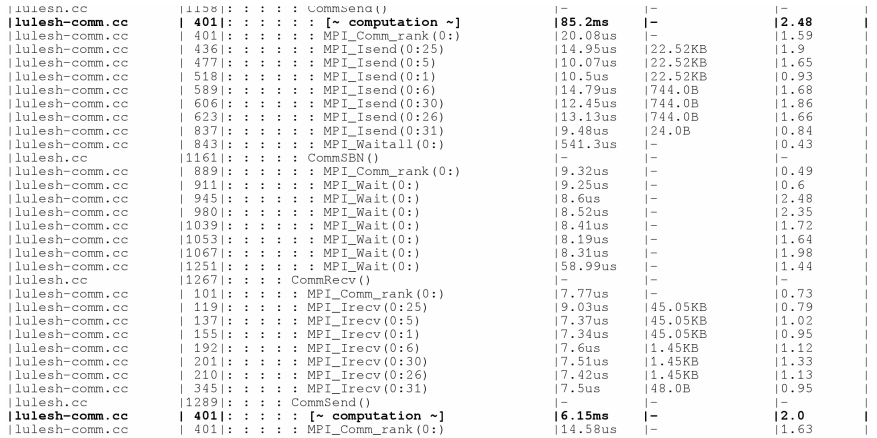

```
|lulesh.cc          |1158|: : : : : CommSend()                    |-          |-          |-          |
|lulesh-comm.cc     | 401|: : : : : : [~ computation ~]           |85.2ms     |-          |2.48       |
|lulesh-comm.cc     | 401|: : : : : : MPI_Comm_rank(0:)           |20.08us    |-          |1.59       |
|lulesh-comm.cc     | 436|: : : : : : MPI_Isend(0:25)             |14.95us    |22.52KB    |1.9        |
|lulesh-comm.cc     | 477|: : : : : : MPI_Isend(0:5)              |10.07us    |22.52KB    |1.65       |
|lulesh-comm.cc     | 518|: : : : : : MPI_Isend(0:1)              |10.5us     |22.52KB    |0.93       |
|lulesh-comm.cc     | 589|: : : : : : MPI_Isend(0:6)              |14.79us    |744.0B     |1.68       |
|lulesh-comm.cc     | 606|: : : : : : MPI_Isend(0:30)             |12.45us    |744.0B     |1.86       |
|lulesh-comm.cc     | 623|: : : : : : MPI_Isend(0:26)             |13.13us    |744.0B     |1.66       |
|lulesh-comm.cc     | 837|: : : : : : MPI_Isend(0:31)             |9.48us     |24.0B      |0.84       |
|lulesh-comm.cc     | 843|: : : : : : MPI_Waitall(0:)             |541.3us    |-          |0.43       |
|lulesh.cc          |1161|: : : : : CommSBN()                     |-          |-          |-          |
|lulesh-comm.cc     | 889|: : : : : : MPI_Comm_rank(0:)           |9.32us     |-          |0.49       |
|lulesh-comm.cc     | 911|: : : : : : MPI_Wait(0:)                |9.25us     |-          |0.6        |
|lulesh-comm.cc     | 945|: : : : : : MPI_Wait(0:)                |8.6us      |-          |2.48       |
|lulesh-comm.cc     | 980|: : : : : : MPI_Wait(0:)                |8.52us     |-          |2.35       |
|lulesh-comm.cc     |1039|: : : : : : MPI_Wait(0:)                |8.41us     |-          |1.72       |
|lulesh-comm.cc     |1053|: : : : : : MPI_Wait(0:)                |8.19us     |-          |1.64       |
|lulesh-comm.cc     |1067|: : : : : : MPI_Wait(0:)                |8.31us     |-          |1.98       |
|lulesh-comm.cc     |1251|: : : : : : MPI_Wait(0:)                |58.99us    |-          |1.44       |
|lulesh.cc          |1267|: : : : CommRecv()                      |-          |-          |-          |
|lulesh-comm.cc     | 101|: : : : : MPI_Comm_rank(0:)             |7.77us     |-          |0.73       |
|lulesh-comm.cc     | 119|: : : : : MPI_Irecv(0:25)               |9.03us     |45.05KB    |0.79       |
|lulesh-comm.cc     | 137|: : : : : MPI_Irecv(0:5)                |7.37us     |45.05KB    |1.02       |
|lulesh-comm.cc     | 155|: : : : : MPI_Irecv(0:1)                |7.34us     |45.05KB    |0.95       |
|lulesh-comm.cc     | 192|: : : : : MPI_Irecv(0:6)                |7.6us      |1.45KB     |1.12       |
|lulesh-comm.cc     | 201|: : : : : MPI_Irecv(0:30)               |7.51us     |1.45KB     |1.33       |
|lulesh-comm.cc     | 210|: : : : : MPI_Irecv(0:26)               |7.42us     |1.45KB     |1.13       |
|lulesh-comm.cc     | 345|: : : : : MPI_Irecv(0:31)               |7.5us      |48.0B      |0.95       |
|lulesh.cc          |1289|: : : : CommSend()                      |-          |-          |-          |
|lulesh-comm.cc     | 401|: : : : : [~ computation ~]             |6.15ms     |-          |2.0        |
|lulesh-comm.cc     | 401|: : : : : MPI_Comm_rank(0:)             |14.58us    |-          |1.63       |
```

Table: Total MPI calls

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Rank 0	300	510	510	90	29	270

Table: MPI calls counts by communication phases

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Comm phase 1	0	7	0	0	1	1
Comm phase 2	7	7	7	1	0	3
Comm phase 3	0	0	7	1	0	2
Comm phase 4	0	3	0	0	0	1
Comm phase 5	3	0	3	1	0	2

Validation

Lulesh 2.0

```
|lulesh.cc          |1289|: : : : CommSend()                  |-          |-  |-          |
|lulesh-comm.cc     | 401|: : : : : [~ computation ~]         |6.15ms     |-  |2.0        |
|lulesh-comm.cc     | 401|: : : : : MPI_Comm_rank(0:)         |14.58us    |-  |1.63       |
|lulesh-comm.cc     | 843|: : : : : MPI_Waitall(0:)           |10.26us    |-  |0.35       |
|lulesh.cc          |1292|: : : : CommSyncPosVel()            |-          |-  |-          |
|lulesh-comm.cc     |1310|: : : : : MPI_Comm_rank(0:)         |22.14us    |-  |0.73       |
|lulesh-comm.cc     |1332|: : : : : MPI_Wait(0:)              |274.73us   |-  |0.79       |
|lulesh-comm.cc     |1366|: : : : : MPI_Wait(0:)              |72.95us    |-  |2.46       |
|lulesh-comm.cc     |1402|: : : : : MPI_Wait(0:)              |23.24us    |-  |2.3        |
|lulesh-comm.cc     |1461|: : : : : MPI_Wait(0:)              |198.36us   |-  |1.79       |
|lulesh-comm.cc     |1475|: : : : : MPI_Wait(0:)              |7.83us     |-  |1.83       |
|lulesh-comm.cc     |1489|: : : : : MPI_Wait(0:)              |7.75us     |-  |2.08       |
|lulesh-comm.cc     |1674|: : : : : MPI_Wait(0:)              |333.07us   |-  |1.73       |
|lulesh.cc          |2656|: : : LagrangeElements()            |-          |-  |-          |
|lulesh.cc          |2461|: : : : CalcQForElems()             |-          |-  |-          |
|lulesh.cc          |1992|: : : : : CommRecv()                |-          |-  |-          |
|lulesh-comm.cc     | 101|: : : : : : [~ computation ~]       |29.32ms    |-  |2.25       |
|lulesh-comm.cc     | 101|: : : : : : MPI_Comm_rank(0:)       |18.76us    |-  |1.39       |
```

Table: Total MPI calls

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Rank 0	300	510	510	90	29	270

Table: MPI calls counts by communication phases

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Comm phase 1	0	7	0	0	1	1
Comm phase 2	7	7	7	1	0	3
Comm phase 3	0	0	7	1	0	2
Comm phase 4	0	3	0	0	0	1
Comm phase 5	3	0	3	1	0	2

Validation

Lulesh 2.0

```
|lulesh.cc          |1992|: : : : : CommRecv()                    |          |               |          |
|lulesh-comm.cc     | 101|: : : : : : [~ computation ~]           |29.32ms   |-              |2.25      |
|lulesh-comm.cc     | 101|: : : : : : MPI_Comm_rank(0:)           |18.76us   |-              |1.39      |
|lulesh-comm.cc     | 119|: : : : : : MPI_Irecv(0:25)             |12.61us   |21.09KB        |0.43      |
|lulesh-comm.cc     | 137|: : : : : : MPI_Irecv(0:5)              |9.23us    |21.09KB        |1.01      |
|lulesh-comm.cc     | 155|: : : : : : MPI_Irecv(0:1)              |8.79us    |21.09KB        |1.17      |
|lulesh.cc          |2010|: : : : : CommSend()                    |-         |-              |-         |
|lulesh-comm.cc     | 401|: : : : : : [~ computation ~]           |15.35ms   |-              |1.91      |
|lulesh-comm.cc     | 401|: : : : : : MPI_Comm_rank(0:)           |15.24us   |-              |1.52      |
|lulesh-comm.cc     | 436|: : : : : : MPI_Isend(0:25)             |13.14us   |21.09KB        |1.9       |
|lulesh-comm.cc     | 477|: : : : : : MPI_Isend(0:5)              |10.28us   |21.09KB        |1.89      |
|lulesh-comm.cc     | 518|: : : : : : MPI_Isend(0:1)              |10.57us   |21.09KB        |1.26      |
|lulesh-comm.cc     | 843|: : : : : : MPI_Waitall(0:)             |291.68us  |-              |0.45      |
|lulesh.cc          |2014|: : : : : CommMonoQ()                   |-         |-              |-         |
|lulesh-comm.cc     |1734|: : : : : : MPI_Comm_rank(0:)           |9.51us    |-              |0.45      |
|lulesh-comm.cc     |1757|: : : : : : MPI_Wait(0:)                |8.88us    |-              |0.81      |
|lulesh-comm.cc     |1791|: : : : : : MPI_Wait(0:)                |8.78us    |-              |2.04      |
|lulesh-comm.cc     |1824|: : : : : : MPI_Wait(0:)                |8.53us    |-              |2.13      |
|                   |    |: END LOOP                              |-         |-              |-         |
+-------------------+----+----------------------------------------+----------+---------------+----------+
```

Table: Total MPI calls

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Rank 0	300	510	510	90	29	270

Table: MPI calls counts by communication phases

	MPI_Isend	MPI_Irecv	MPI_Wait	MPI_Waitall	MPI_Allreduce	MPI_Comm_rank
Comm phase 1	0	7	0	0	1	1
Comm phase 2	7	7	7	1	0	3
Comm phase 3	0	0	7	1	0	2
Comm phase 4	0	3	0	0	0	1
Comm phase 5	3	0	3	1	0	2

Outline

Scalability

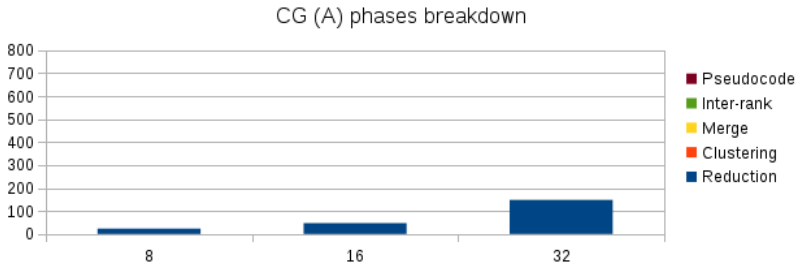

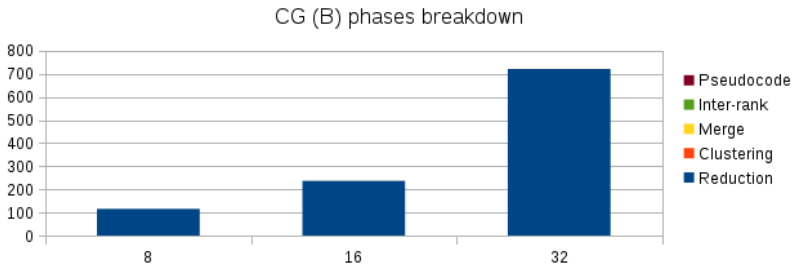

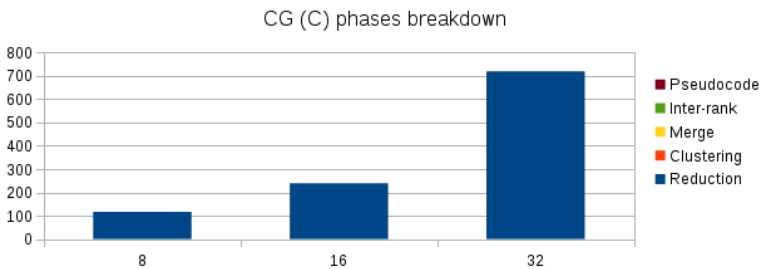

Scalability

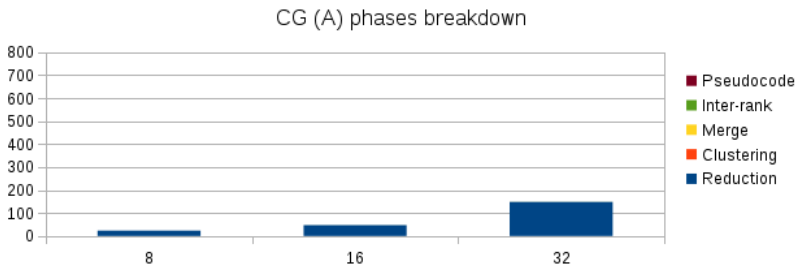

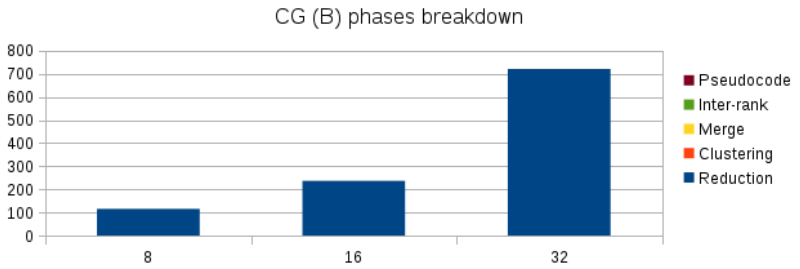

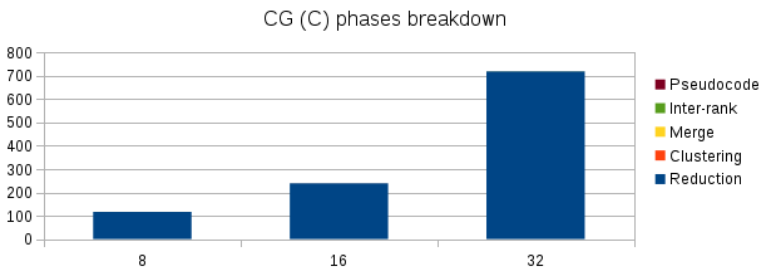

Reduction step is clearly dominating the execution time.

Scalability

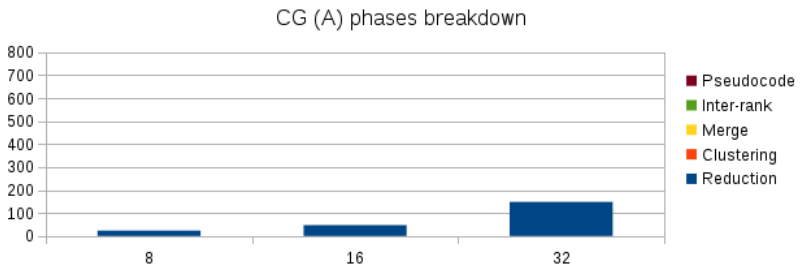

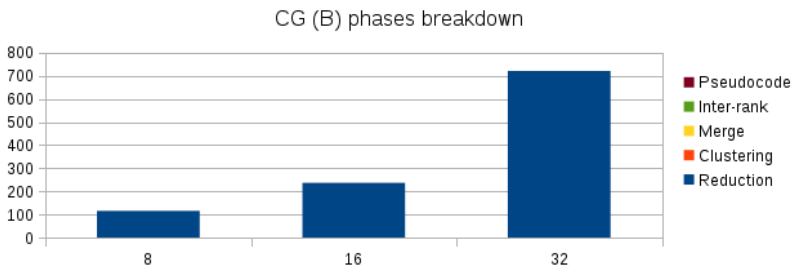

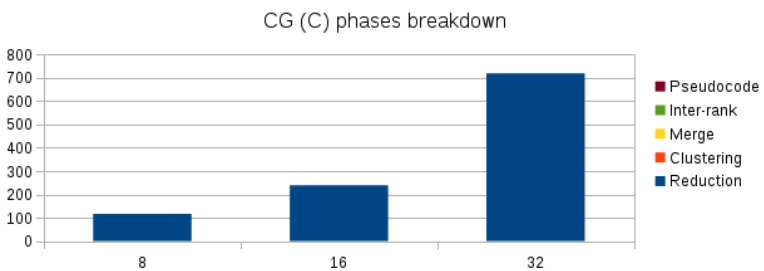

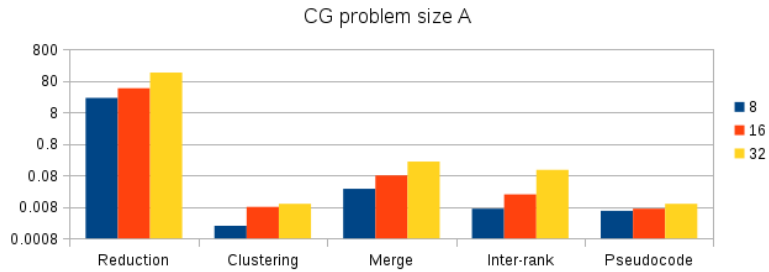

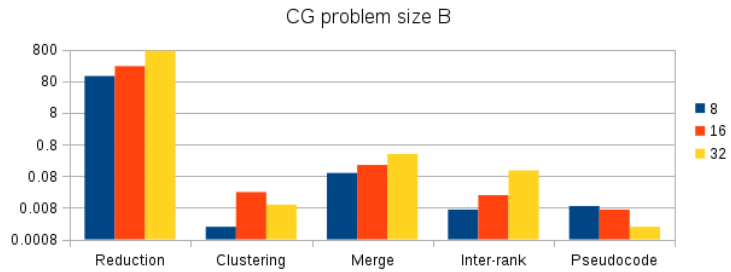

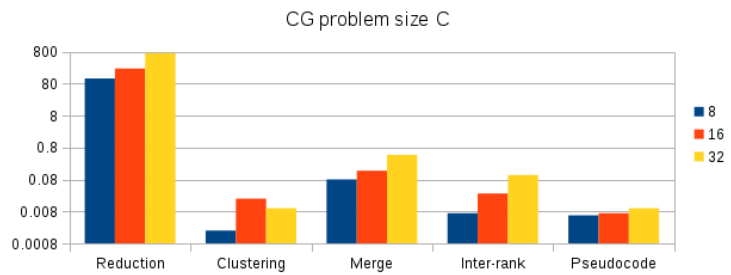

Reduction step is clearly dominating the execution time.

Scalability

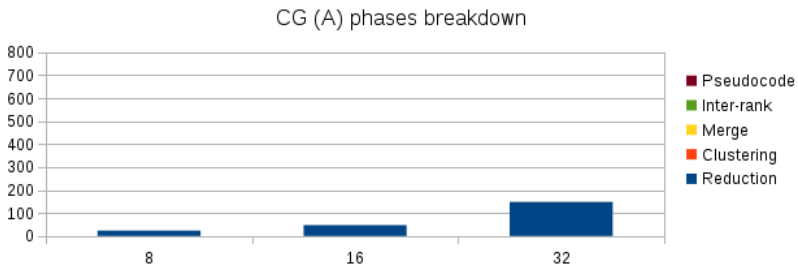

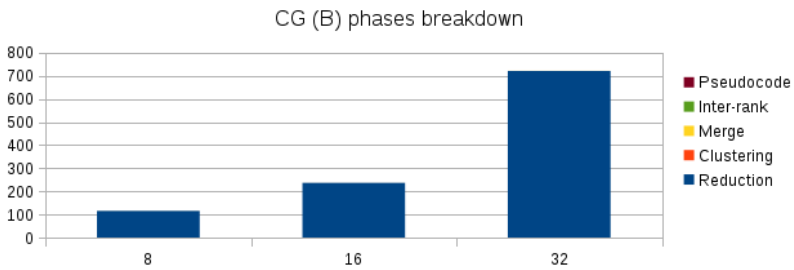

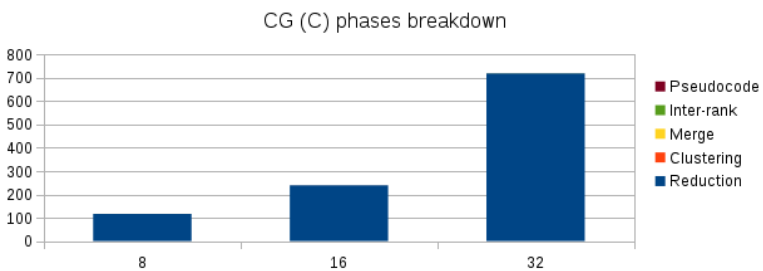

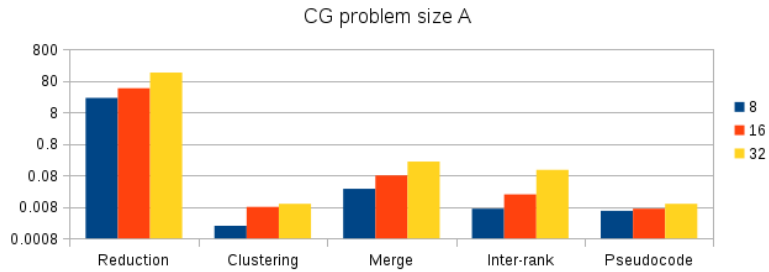

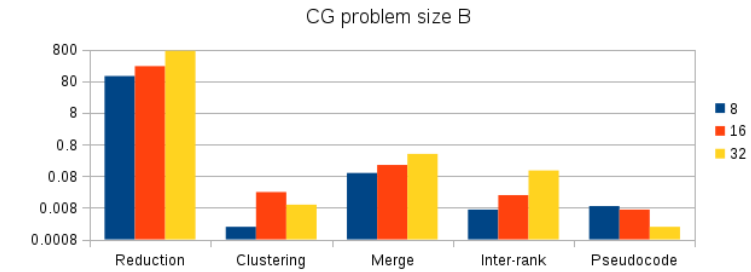

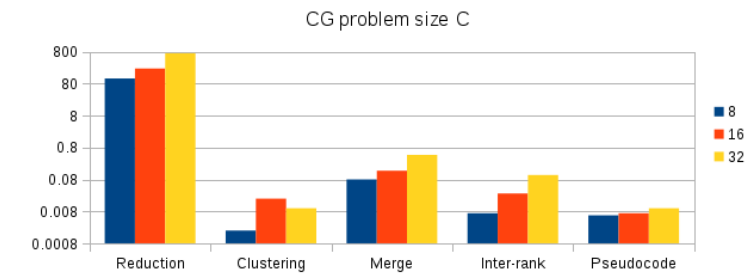

Reduction step is clearly dominating the execution time.

The rest maintains its time even if increasing problem size.

Scalability

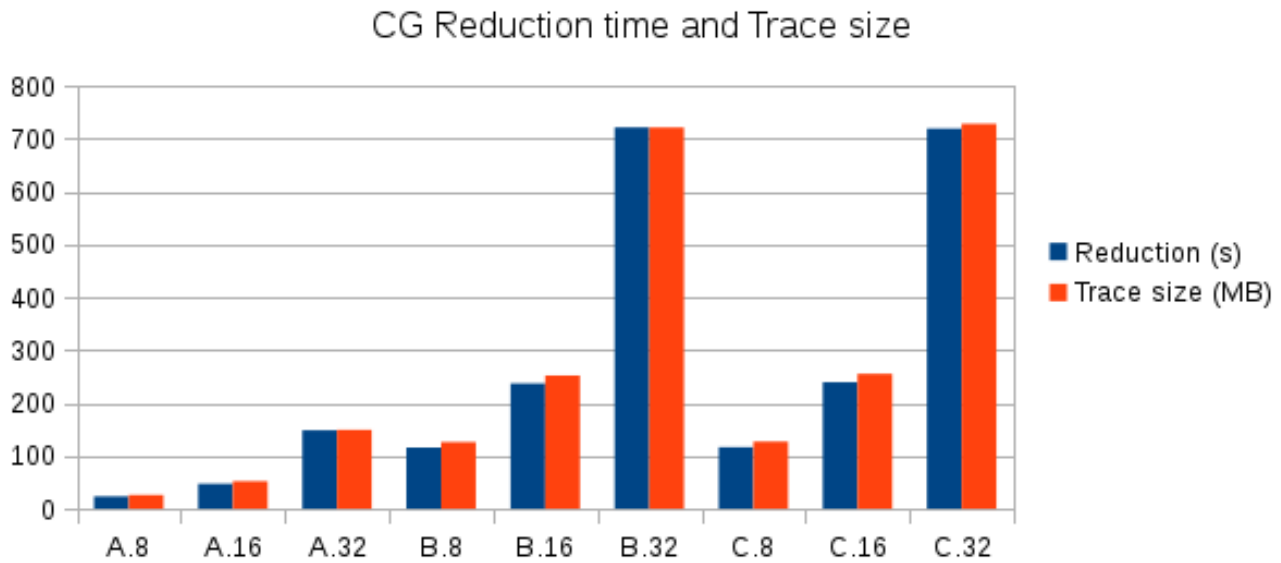

Keynote

Reduction time is highly correlated with trace size